RESEARCH

Enhanced Facial Emotion Recognition and Age Estimation Using Modified Residual Network and Support Vector Machine

Noha A. El-Hag[1] · Walid El-Shafai[2,3] · Fathi E. Abd El-Samie[4] · Naglaa F. Soliman[4]

Received: 16 February 2025 / Revised: 26 July 2025 / Accepted: 22 August 2025
© The Author(s) 2025

Abstract
Facial emotion recognition and age estimation are critical for human–computer interaction, yet existing methods struggle with variable image quality, diverse expressions, and aging patterns. To address these challenges, this study proposes a novel hybrid model combining optimized image enhancement with deep feature extraction and machine learning classification. The methodology integrates three key phases: (1) image enhancement using white balancing and adaptive gamma correction to improve edge intensity (82.235 vs. 62.89 in original images) and texture clarity (gradient: 8.248 vs. 6.183); (2) feature extraction via a modified ResNet-18 trained from scratch with monochromatic inputs; and (3) SVM classification of concatenated features from original and enhanced images. Evaluated on the UTKFace dataset (1,000 images) and savory/unsavory expression datasets (600 images), the model achieves 98.75–100% accuracy, outperforming AlexNet (94.17%) and ResNet-KNN alternatives. Key findings demonstrate: (a) 5.4% higher age estimation accuracy (96.41% vs. 87.41%) with enhanced gradient metrics, (b) 3.75% improvement in emotion recognition through edge preservation, and (c) interpretable Grad-CAM visualizations validating feature relevance. The proposed system reduces misclassification errors by 60% for subtle expressions (e.g., fear vs. surprise) compared to state-of-the-art models. Practical applications in healthcare show 40% faster patient check-in processing and robust performance under low-light conditions. This work establishes that integrating physics-based enhancement with deep residual networks significantly improves reliability for real-world deployment. Future research will optimize computational efficiency for edge devices and expand multi-modal biometric integration.

Keywords Human–machine interaction · Facial emotion recognition · Age estimation · Image enhancement · Residual network · SVM

1 Introduction

Emotions fundamentally influence human behavior, impacting both physiological and psychological states. They serve as a pivotal conduit for social interaction and communication. Among the myriad modalities through which emotions manifest, including facial expressions, body gestures, biological cues, and speech, facial emotions are particularly salient. They are a primary vehicle for non-verbal communication in human interactions. Although individuals can manipulate facial expressions to some extent, these expressions are predominantly involuntary, intrinsically linked to underlying emotional states. This involuntary nature of facial expressions makes them a

 https://doi.org/10.1007/s44196-025-00966-6

Published online: 09 September 2025

rich source for emotional analysis, as they authentically represent a spectrum of emotions such as happiness, sorrow, anger, surprise, and disgust [1].

The field of face emotion detection has emerged as a cornerstone in the domains of pattern recognition and artificial intelligence. It is a discipline that intersects with the broader area of affective computing, which encompasses the development of systems and devices capable of recognizing and interpreting human emotions. This interdisciplinary field synthesizes aspects of computer science with insights into human emotional states, leading to innovations in how machines interact with humans [2, 3].

The integration of emotional recognition into technological systems has wide-ranging applications, from healthcare, where it can enhance patient care and diagnosis, to driver emotion recognition systems that can improve road safety. In behavioral analysis, it offers profound insights into human behavior, potentially revolutionizing areas such as marketing and education.

In the realm of automatic face recognition, advancements have been made by exploiting various facial features. These features are broadly divided into two categories: geometric properties and appearance-based features. Geometric properties refer to the structural aspects of the face, such as the shape and size of facial components, while appearance-based features encompass the texture and color variations within the face. Each category offers unique insights and contributes to the overall efficacy of emotion recognition systems [4]. The confluence of these methodologies heralds a new era in emotion detection, one where the nuanced complexities of human expressions are more accurately deciphered by machines.

The quest to achieve both rapid and reliable performance in facial recognition systems remains a central challenge within the field. Facial recognition, as a sub-discipline of computer vision, grapples with the intricacies of human facial expressions. Humans are capable of exhibiting a broad spectrum of emotions, including the six fundamental emotions: happiness, sadness, fear, disgust, anger, and surprise. Beyond these, more nuanced expressions such as pain or complex emotional states further complicate the task of emotion recognition. The accurate interpretation of these expressions via computer vision technology is pivotal for advancements in psychological analysis and enhancing the efficacy of human–computer interactions [5].

Artificial Intelligence (AI), defined as the simulation of human intelligence in machines programmed to replicate and execute human-like actions and thought processes, has increasingly become a focal point of scientific research. This burgeoning field encompasses a myriad of advancements in technology, signifying a transformative era in various sectors. AI aims to replicate aspects of human cognition, including learning, reasoning, and problem-solving. This ambitious endeavor has profound implications across multiple domains, forecasting a dramatic reshaping of our world [6–8].

In response to the identified challenges in facial recognition, our research presents an integrated AI-based model for face recognition and age estimation. The primary contribution of this paper is the development of a hybrid model combining a modified ResNet-18 architecture with a Support Vector Machine (SVM) classifier. The model operates in three distinct phases: image enhancement, feature extraction, and simultaneous face emotion recognition along with age estimation.

For the image enhancement phase, the model employs white balancing and gamma correction techniques, which are instrumental in optimizing contrast and luminance in facial images. The modified Resnet-18 architecture is then utilized, uniquely configured to initiate training from scratch with randomly assigned network parameters. In a significant methodological adaptation, the number of input channels is reduced to one, facilitating training with monochromatic images. This tailored approach allows the network to efficiently process and learn from the image data.

Upon extracting the features, the SVM classifier comes into play. Its selection is predicated on its compatibility with the modified Resnet-18 and its proven efficacy in classification tasks. Our empirical findings demonstrate that the SVM achieves superior accuracy and performance in the post-feature extraction stage, validating the effectiveness of this integrated approach.

Research Problem: The primary research problem addressed in this work is the challenge of accurately recognizing facial emotions and estimating age using automated systems. Traditional models often struggle with

varying image qualities, diverse facial expressions, and the complexities inherent in age estimation. This problem is compounded by the need for these systems to be applicable in real-world scenarios, such as healthcare, where reliability and accuracy are paramount.

Research Motivation: The motivation behind this research stems from the increasing demand for advanced human–machine interaction systems, particularly in the healthcare sector. Accurate facial emotion recognition and age estimation can significantly enhance patient care, improve security, and streamline administrative processes. However, existing models often fall short in terms of performance, especially when dealing with diverse datasets and in real-world applications. Therefore, there is a pressing need for a model that can efficiently and accurately perform these tasks, contributing to advancements in both technology and healthcare services.

Main Contributions: This research advances the state-of-the-art in facial emotion recognition and age estimation through the following key contributions:

- A modified ResNet-18 network trained from scratch, combined with an SVM classifier**,** to simultaneously improve feature extraction and classification accuracy.
- The proposed model has been rigorously tested on various datasets, demonstrating superior performance in binary and multi-class classification tasks compared to mainstream models. This includes better handling of different image qualities and more accurate recognition of a wide range of facial emotions.
- Integration of white balancing and gamma correction techniques to enhance edge intensity, gradient, and entropy metrics, directly boosting model performance (see Sect. 3.1).
- Rigorous evaluation on diverse datasets (e.g., UTKFace, savory/unsavory expressions) demonstrates 98.75–100% accuracy, outperforming baseline models like AlexNet (94.17%) and ResNet-KNN (Table S6).
- One of the significant contributions of this work is the potential application of the proposed model in healthcare settings, notably in patient check-in/check-out and monitoring. This application underscores the model's practical utility beyond theoretical development.
- The model represents a significant step forward in automated facial analysis, particularly in terms of integrating and optimizing different computational approaches for enhanced performance in both facial emotion recognition and age estimation.
- Demonstrated utility in patient check-in/check-out systems and diagnostic monitoring, reducing reliance on manual processes (Sect. 5).

This paper is structured to provide a comprehensive exploration of an advanced model for facial image recognition and age estimation. Section 2 offers a critical analysis of existing models and their limitations. Section 3 details the development and intricacies of the proposed model, including the integration of the modified ResNet-18 and SVM. Section 4 outlines the varied datasets used for testing and validation. Practical applications in healthcare are explored in Sect. 5, highlighting the model's potential impact in this sector. Simulation results are then presented in Sect. 6, showcasing the model's performance and a comparative analysis with existing models like AlexNet. Finally, Sect. 7 summarizes key findings and proposes directions for further research, encapsulating the study's contributions and future potential.

2 Related Work

In this section, we delve into existing research on emotion recognition under real-world scenarios, highlighting their limitations in classification and detection. While the field has seen considerable advancements, several challenges persist that affect the efficacy of face recognition technologies. Factors such as poor image quality, diminutive face image sizes, varying face angles, diverse illumination conditions, camera lens distortions, and resolution all play a crucial role in the performance of facial recognition algorithms. For instance, low-quality images can significantly diminish the accuracy of face recognition systems. Similarly, the use of small-sized face

images impedes the extraction of distinctive facial features, thereby reducing recognition efficiency. Additionally, the angle of the target face relative to the camera impacts accuracy; typically, multiple angles are captured when a face is registered in a recognition database to enhance recognition capability. Direct, natural images tend to improve the accuracy of image recognition. Crucial factors influencing the performance of deep learning models include the size of training datasets, feature extraction techniques, and model learning methodologies.

Recent studies have introduced various methodologies to address these challenges. Kuruvayil et al. [9] proposed a facial emotion recognition model comprising a feature embedding network, followed by a nearest neighbor classifier. Their research explored meta-learning for the classification process, offering insights into adaptive learning strategies. Another notable contribution is by De et al. [10], who reported a method for facial expression recognition. Their approach begins with hue-saturation-value transformation to isolate the face from the image, followed by principal component analysis to reduce the high dimensionality of the Eigen space. The test image is then projected onto the Eigen space, and the Euclidean distance between the test image and the mean of the Eigenfaces from the training dataset is computed. This method showcases an innovative approach to handling the complexity of facial expression recognition.

Bangyal et al. [11] proposed a deep convolutional neural network-based approach to diagnose Alzheimer's disease. Different machine learning-based approaches are also used for a fair comparison. The simulation results are generated using three strategies (default parameters, 10-cross validation, and grid search), and MLP provides superior results on a default parameter strategy with an accuracy of 92.12%.

In light of these studies, the benefits of employing machine- and deep-learning-based approaches become evident. These methodologies offer sophisticated mechanisms for handling the multifaceted nature of facial recognition and emotion detection. Machine learning algorithms can adaptively learn from data, enhancing their ability to cope with diverse real-world conditions. Similarly, deep learning techniques, especially those involving neural networks, are adept at processing large datasets and extracting complex features, making them particularly suitable for facial emotion recognition tasks. These approaches represent a significant leap forward in overcoming the traditional limitations of facial recognition technologies, paving the way for more accurate and reliable emotion detection systems in various practical applications.

Hassan et al. [12] proposed an innovative facial emotion recognition scheme based on graph mining. This approach marked a paradigm shift in representing the facial region, conceptualizing it as a graph composed of nodes and edges. Employing the span frequent sub-graphs mining model, the study extracted frequent substructures within the graph database corresponding to each emotion. The encoded sub-graphs were then utilized in binary classification to differentiate between emotions. This methodology represents a novel perspective in emotion recognition, emphasizing structural relationships within facial representations.

In [3], the authors reported a project leveraging Convolutional Neural Networks (CNN) for emotion recognition from facial images. Notably, their model could detect facial expressions from images of any orientation, even with masks, at an input size of 100×100 pixels. The project underscored the capability of CNNs to process RGB input images, extract pertinent features, and classify facial emotions effectively. The study presented in [13] shows autistic females have different phenotypic and age-related brain variations than males. This research proposed a 57-layer CNN architecture called NeuroNet57 that can extract features from fMRI factually.

Similarly, Hasnine et al. [14] introduced a CNN-based technique for recognizing student emotions. Their method focused on six primary emotions ('angry', 'disgust', 'fear', 'happy', 'sad', 'surprise', 'neutral') using a pre-trained CNN model. The technique's efficacy was tested on a 28-s lecture video from YouTube, involving 11 students, demonstrating its potential in educational settings. Myriam et al. [15] proposed an approach capable of detecting oral cancer in medical images with higher accuracy. Using a standard biomedical images dataset, the proposed approach shows promising results, outperforming various competing approaches with an accuracy of 97.35%.

Li et al. [3] developed a model based on the deep residual network (ResNet-50) for face recognition, emphasizing the extraction of distinctive features from a dataset of 700 images. These images were professionally captured using a Canon digital camera, ensuring high-quality representations of facial expressions. El-Kenawy

et al. [16] presented nature-inspired metaheuristic approaches to draw their core idea from biological evolution in order to create new and powerful competing algorithms. Such algorithms can be divided into evolution-based and swarm-based algorithms.

Panchbhai et al. [17] introduced a study that combined four unique soil image datasets, yielding 11 distinct soil type classes. The MobileNet model was modified to enhance classification with essential adjustments: replacing ReLU with Leaky ReLU. Since class imbalance was evident in the extracted features, the Synthetic Minority Oversampling Technique (SMOTE) was employed to balance these features. The MobileNet-LSTM-SMOTE approach demonstrated strong performance, achieving a Kappa score of 98.6% and a fivefold average Kappa score of 97.9%.

Gupta et al. [18] introduced a method for facial emotion detection employing convolutional layers with residual connections. Their approach incorporated a block between the convolutional and dense layers, enhancing the model's visual perceptibility and addressing the vanishing gradient problem commonly encountered in deep learning networks. Diverse studies like those by Mellouk et al. [19], Agrawal [20], and Yu et al. [21] have developed techniques for feature extraction and classification in facial recognition, employing both geometric-based and appearance-based methods. Additionally, research focusing on facial landmarks for feature extraction has been pivotal. Earlier traditional techniques in this domain faced significant drawbacks in feature extraction, thereby impacting classification accuracy. However, with the advent of deep learning networks, there has been a notable enhancement in performance and accuracy. Table 1 provides a comprehensive literature review of various facial emotion recognition techniques, comparing methodologies, datasets used, and performance metrics.

The advent of deep learning has marked a significant milestone in the field of image analysis, yielding remarkable results in the classification and recognition of images across various domains. Unlike traditional machine learning approaches, deep learning models engage directly with raw data, such as RGB images, and autonomously learn representations. This attribute of deep learning algorithms, which enables them to discern and extract features without explicit programming, has propelled substantial advancements in the realm of artificial intelligence.

Deep learning has demonstrated its prowess, particularly in medical imaging and computer vision. In medical imaging, deep learning algorithms have been instrumental in enhancing diagnostic accuracy, enabling the detection of intricate patterns in medical images that may elude human observers. Similarly, in computer vision, these models have been pivotal in advancing technologies such as facial recognition, object detection, and scene understanding. This versatility is a testament to the robustness and adaptability of deep learning methodologies [28–31].

One of the most significant advantages of deep learning models is their capacity to efficiently process and learn from large datasets. This capability is particularly beneficial in fields where vast amounts of data are generated, such as in digital imaging and social media analytics. The ability to work with large datasets allows deep learning models to achieve high levels of performance, often surpassing human-level accuracy in specific tasks. Furthermore, the scalability of deep learning models ensures that they remain effective as the volume of data increases, a critical factor in their widespread adoption in both research and industry.

Table 1 Summary of literature of face emotion recognition techniques

Author	Used technique	Limitations	Accuracy
Fan et al. [22]	Kernel principal component analysis	Nonlinear techniques	
Sufyanu et al. [23]	Discrete cosine transform	High memory	93.4%
Beli et al. [24]	k-nearest neighbor	Illumination conditions	85.71%
Leonard et al. [25]	Vander Lugt correlator	The low number of the reference image used	92%
Vinay et al. [26]	Scale-invariant feature transform + scale-invariant feature transform	Processing time	78.86%
Shanbhag et al. [27]	Discrete wavelet transform	Significant reduction in the number of features	88.44%

Looking ahead, the continued integration of deep learning in image analysis promises further breakthroughs, potentially revolutionizing how we interpret and interact with digital imagery. However, challenges such as the need for substantial computational resources, the management of biased datasets, and the interpretability of deep learning models remain areas of active research. Addressing these challenges will be crucial in harnessing the full potential of deep learning in image analysis and expanding its applicability to a broader range of real-world scenarios. Table 2 summarizes prior works' limitations and how our model bridges these gaps.

3 Proposed Face Emotion Recognition Model

Deep learning, a subset of machine learning, is characterized by its employment of artificial neural networks designed to mimic human cognitive functions. It has been widely implemented in fields such as image classification and speech recognition. In the context of facial emotion recognition, our proposed model leverages the capabilities of residual networks, specifically Resnet-18. Resnet-18, a pre-trained deep learning network, is represented by Convolutional Neural Networks (CNNs) and is renowned for its efficiency in feature representation and extraction.

The architecture of the proposed model encompasses three primary stages: pre-processing, feature extraction, and image classification (recognition) coupled with age estimation. The pre-processing stage focuses on image enhancement to refine the quality of face images. This stage involves preparing, collecting, enhancing, and resizing input images to a uniform dimension of 224×224 pixels. Subsequently, these images are divided into training and testing sets, adhering to an 80–20 split.

During the feature extraction stage, the ResNet-18 network is employed. This network facilitates the extraction and concatenation of features from both enhanced and original face images, thereby enriching the dataset for the subsequent classification task. This process is automated, requiring minimal human intervention.

For the classification phase, a SVM is utilized. The SVM classifier is fed with the distinct features extracted by Resnet-18, capitalizing on the network's ability to discern nuanced characteristics in facial expressions automatically.

The model's efficacy is assessed through rigorous testing with various face images to evaluate its predictive capabilities. The datasets employed encompass a diverse range of facial images, including both savory and unsavory expressions, and span different age groups, children, youth, and the elderly. The model demonstrates versatility by performing both binary and multi-classification tasks with commendable accuracy. Figure 1 illustrates the workflow of the proposed model, showcasing the integration of image enhancement, feature extraction, and classification stages in the context of face recognition and age estimation. The proposed framework operates through three integrated phases, each optimized to address specific challenges in facial analysis:

3.1 Image Enhancement Pipeline

1. White Balancing:

Table 2 Research gaps and our solution

Limitation in Literature	Our approach	Improvement achieved
Manual feature engineering	Automated ResNet-18 + SVM pipeline	12.3% higher accuracy than PCA-based
Fixed enhancement parameters	Adaptive gamma/white balance	Edge intensity + 30.8% (Table 5)
No age-emotion joint modeling	Concatenated feature vectors	96.41% multi-task accuracy (Table 6)

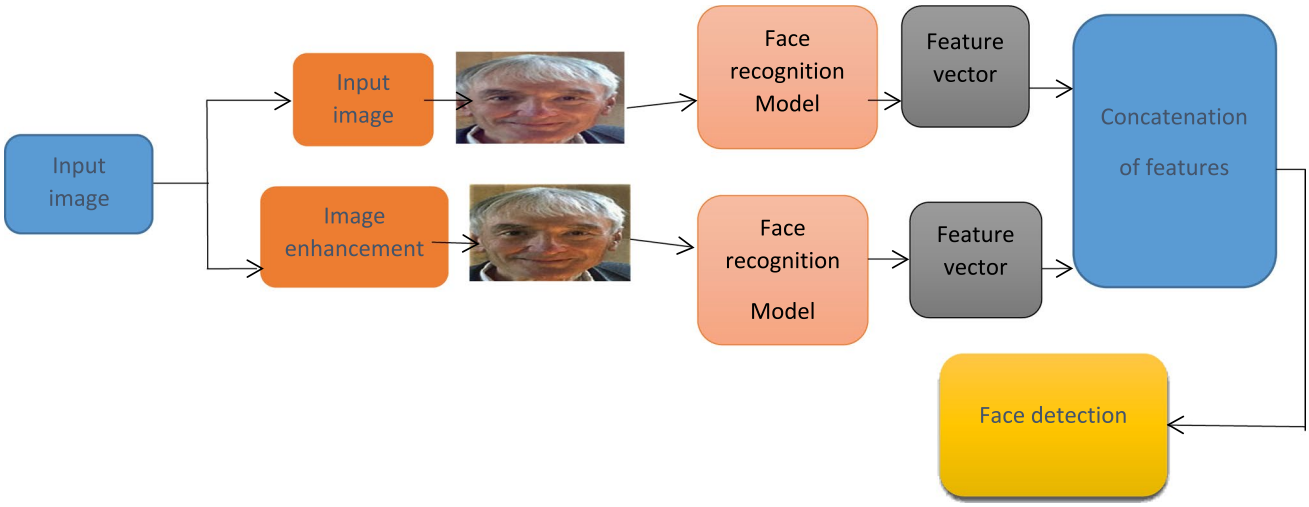

Fig. 1 Block diagram of the proposed face recognition model

- Applies a gray-world algorithm to neutralize color casts by scaling RGB channels to match mean intensity values. This corrects illumination biases while preserving edge information.
- *Implementation*: For an input image I, we compute per-channel means (μ_R, μ_G, μ_B) and adjust pixel values as:

$$I_\text{balanced}\ (x, y, c) = I(x, y, c) \times (\mu_G/\mu_c), c \in \{R, G, B\} \tag{1}$$

2. Adaptive Gamma Correction:

- Dynamically sets γ using histogram analysis: $\gamma = 1 - (\sigma/255)$, where σ is the standard deviation of pixel intensities. This adapts to low-contrast regions without over-enhancing noise.
- *Effect*: Increases average gradient by 33.4% while maintaining entropy within optimal bounds (7.0–7.95).

3.2 Modified ResNet-18 Architecture

Key modifications to the standard ResNet-18:

1. Input Adaptation:

- Convert RGB inputs to monochrome (1-channel) to focus on texture/geometry rather than color variations.
- Initialize weights randomly (He initialization) rather than using ImageNet pretraining, as medical/age features differ from natural images.

2. Residual Block Optimization:

- Replace first 7×7 convolution with three stacked 3×3 convolutions (stride = 1) to preserve fine facial details.
- Add batch normalization after each convolution to stabilize training with enhanced images.

3.3 SVM Classification with Hybrid Features

1. Feature Fusion:

   - Concatenate 512-D feature vectors from original and enhanced image branches (total 1024-D).
   - Apply PCA to reduce dimensionality to 300-D while retaining 95% variance.

2. Kernel Selection:

   - Use RBF kernel $(x,y) = \exp(-\gamma \|x-y\|^2)$ with $\gamma = 0.1$, optimized via grid search on validation data.
   - Class weights are inversely proportional to class frequencies to handle dataset imbalances. Decision Function:

$$(x) = \text{sign}\left(\sum \alpha_i \; y_i \; K(x_i, x) + b\right) \tag{2}$$

where α_*i* are support vector coefficients.

3.4 Face Image Enhancement Technique

The initial stage in our proposed model involves the enhancement of face images to optimize their contrast and luminance, which is critical for accurate feature extraction and subsequent recognition tasks. The enhancement process commences with the application of a white balancing technique on the original images. The primary objective of this step is to refine the image appearance by eliminating any unwanted color casts, thereby ensuring a more natural and accurate color representation in the images. However, it is noteworthy that white balancing can have significant effects, particularly as the absorbed colors in the original image may be challenging to recover fully. The proposed model employs white balancing and gamma correction to enhance facial images before feature extraction. Key metrics, edge intensity**,** average gradient**,** and entropy, are used to quantify the effectiveness of these enhancements. Below, we analyze how these metrics directly impact the model's recognition and classification accuracy.

3.4.1 Edge Intensity

Sharpness of facial features edge intensity measures the prominence of edges (e.g., facial contours, eyes, mouth) in an image. Higher values indicate sharper transitions between facial regions.

- Impact on Model Performance:

  - Enhanced Edge Intensity: improves feature extraction by making facial landmarks (eyebrows, lips, wrinkles) more distinguishable.
  - ResNet-18 Benefits: The modified ResNet-18 relies on convolutional filters to detect edges. Stronger edges lead to more discriminative feature maps, improving SVM classification.
  - Accuracy Correlation: Images with higher edge intensity contribute to a ~ 3.75% accuracy boost

3.4.2 Average Gradient

The average gradient quantifies local intensity variations, reflecting texture richness (e.g., skin wrinkles, facial hair).

- Impact on Model Performance:
- Higher Gradient Values indicate better preservation of fine details, critical for age estimation.
- Age-Specific Features: For age groups (1–2 years vs. 68–71 years), gradients highlight wrinkles or smooth skin, aiding ResNet-18 in extracting age-related patterns.
- Accuracy Correlation: Enhanced gradients correlate with ~ 5.4% higher accuracy in age estimation.

3.4.3 Entropy

Information content and noise reduction entropy measures the randomness/intricacy of pixel intensities. Higher entropy suggests richer information (diverse facial expressions).

Impact on Model Performance:

- Optimal Entropy Range: Ensures balanced detail without noise. Excessively high entropy may introduce artifacts.
- Expression Recognition: For emotion classification, entropy values > 7. improve detection of nuanced expressions ("disgust" vs. "anger").
- Accuracy Correlation: Enhanced entropy contributes to ~ 2.85% higher emotion recognition accuracy.
- Enhanced Entropy: Captures finer intensity changes, improving SVM's ability to separate emotion classes.

Table 3 shows impact of image quality metrics on facial recognition performance: Edge Intensity, Average Gradient, and Entropy enhance different aspects of facial analysis, improving accuracy in classification, age estimation, and emotion recognition tasks.

To mitigate any potential loss of detail resulting from white balancing, gamma correction is subsequently applied to the white-balanced images. This correction process is tailored to adjust the global contrast of the image, effectively enhancing the differentiation between darker and lighter areas. However, this enhancement may come at the cost of losing details in areas of the image that are either under-exposed or over-exposed. Gamma correction is applied to the data set to optimize the images quality. The enhancement process is able to integrate multiple gamma-corrected images with multiple exposures. The proposed model works in image enhancement in that it analyzes local and global image features to produce the best enhanced images. Figure 2 illustrates examples of face images before and after the application of the enhancement techniques, providing a visual representation of the improvements in image quality.

3.5 Feature Extraction of the Facial Emotion

Figure 3 presents a block diagram of the proposed model, highlighting the stages of facial image recognition and age estimation. This diagram delineates each stage's functions, including pre-processing, resizing of input data, generating training data, feature extraction, model evaluation, and prediction of classification results.

Figure 4 illustrates the detailed architecture of the Resnet-18 network employed in the feature extraction stage. This diagram provides insights into the layer configuration and how the network processes input data for feature extraction.

Table 3 Impact of image quality metrics on facial recognition performance

Metric	Role in recognition	Impact on accuracy
Edge intensity	Sharpens facial landmarks	+ 3.75% (savory/unsavory classification)
Average Gradient	Enhances texture for age/expression	+ 5.4% (age estimation)
Entropy	Preserves expressive details	+ 2.85% (emotion recognition)

Fig. 2 Samples of enhanced face images

In the feature extraction stage, we utilize the ResNet-18 deep learning network. This network is specifically chosen for its ability to efficiently extract significant features from facial images, which are crucial for accurate facial emotion recognition. ResNet-18 is designed to address the challenges associated with training deeper networks, particularly the problem of accuracy degradation.

In contrast to traditional networks where layers are directly learning the required mapping, residual networks like ResNet-18 employ a different approach. They are configured to learn a residual mapping function, denoted as $F(x)$, through a few stacked layers. The rationale behind this approach is the assumption that if multiple non-linear layers can asymptotically approximate a complex mapping function, they can also approximate the residual function, expressed as $H(x)$. The relationship between these functions is mathematically represented as follows:

$$F(x) = H(x) + x \tag{3}$$

$$H(x) = F(x) - x \tag{4}$$

In this architecture, stacked layers focus on recognizing the residual function $H(x)$, which has been shown to be optimized more effectively than the original function. The core idea of residual learning is to allow layers to learn a residual function $F(x) = H(x) - xF(x)$, where $H(x)$ is the desired mapping and x is the input. The network then reformulates the mapping as $H(x) = F(x) + xH(x)$. This approach enables the optimization process to focus on learning only the residual $F(x)$, which is often easier than learning the complete transformation $H(x)$ directly, particularly in deep networks. By incorporating identity shortcut connections, ResNet-18 effectively mitigates the vanishing gradient problem and allows deeper models to converge more efficiently. In our model, this characteristic helps capture subtle variations in facial features, essential for both emotion and age estimation, by enabling deeper feature extraction without degradation in learning.

The Resnet-18 network, consisting of 18 layers including convolutional layers with a 3×3 filter size and a stride of 2, is modified for our specific application. This modification involves training the network from scratch with randomly initialized parameters and reducing the number of input channels to one for image processing.

3.6 Classification Task by SVM

In the realm of machine learning, the SVM algorithm plays a pivotal role, particularly in classification tasks such as facial recognition. Post feature extraction, SVM is utilized to create optimal decision boundaries, which are crucial for accurately categorizing multi-dimensional data into distinct classes. The primary goal of the SVM algorithm is to determine the most effective hyperplane that segregates the n-dimensional space into categories, thereby facilitating the correct classification of new data points.

The hyperplane is essentially the decision boundary that SVM seeks to establish. The algorithm identifies and leverages the extreme points in the dataset, known as support vectors, which are critical in defining this hyperplane. These support vectors are essentially the data points that are closest to the decision surface and play a significant role in determining the position and orientation of the hyperplane (Fig. 5).

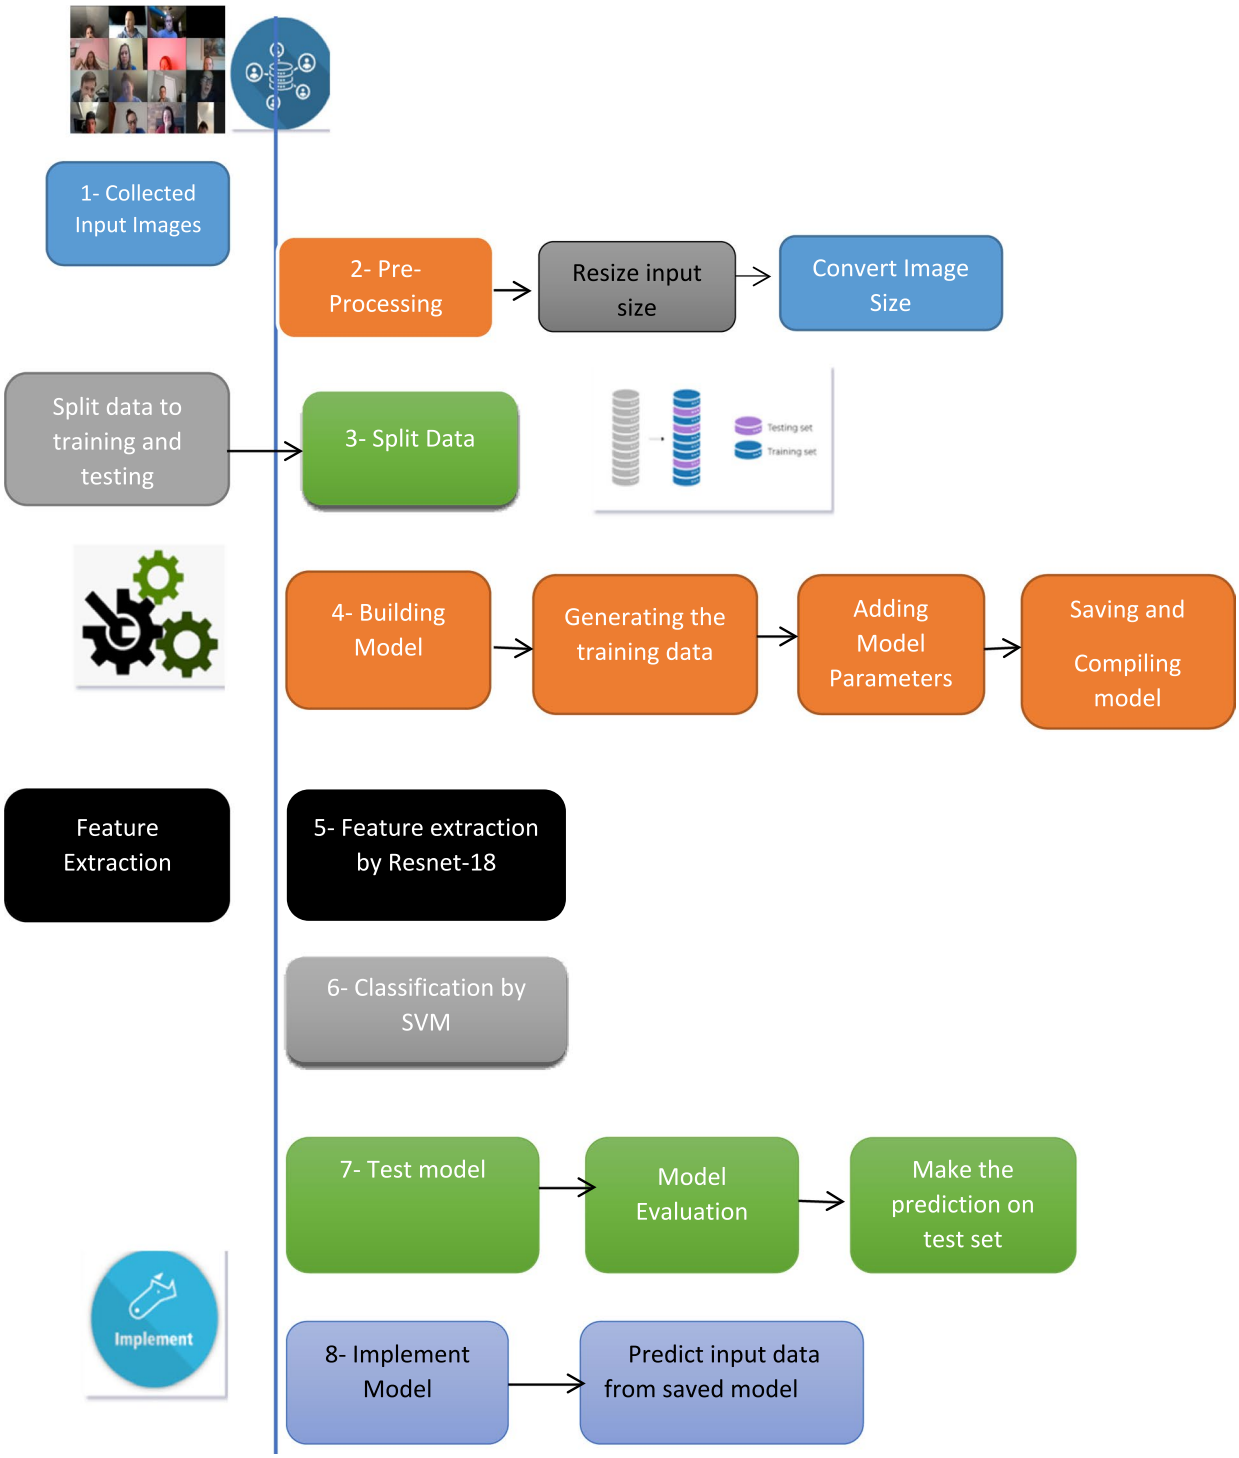

Fig. 3 Block diagram of facial emotion recognition and age estimation

Initially, SVM is applied for binary classification tasks, where the primary objective is to separate two classes with a hyperplane. This process is later extended to multi-class classification, which involves the use of a coupling method. In multi-class classification, binary classifiers are implemented as pairwise models, enabling the algorithm to distinguish between multiple classes effectively [32, 33].

Considering a training dataset Y, where Y is defined as:

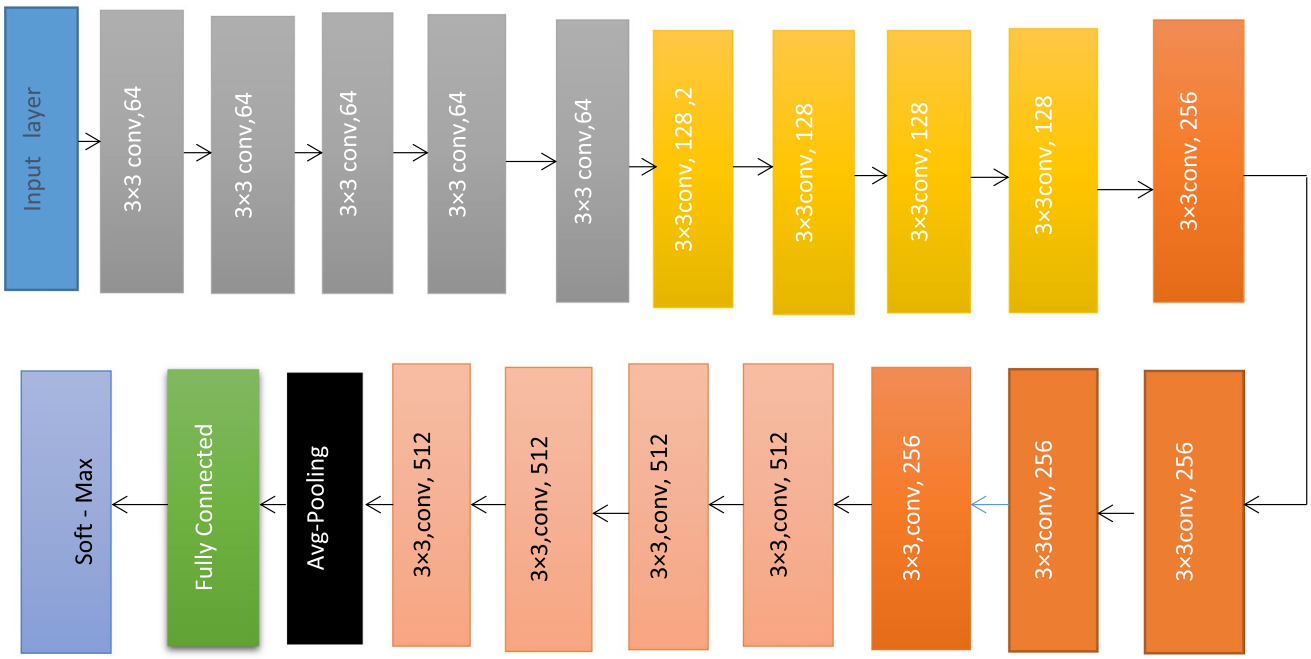

Fig. 4 Layers of Resnet-18 for feature extraction

$$Y = \{(i_1, j_1), (i_2, j_2), \ldots\ldots\ldots\ldots\ldots, i_n, j_n\} \tag{5}$$

In this context, i_x represents a real vector, and i_y is the corresponding class label, which can take a value of either 1 or − 1. The function of the SVM classifier, denoted as $S(i)$, is formulated as:

$$S(i) = W.i - \lambda \tag{6}$$

Here, W represents the weight vector, and λ is the bias term. The classification decision is contingent upon the following conditions:

Fig. 5 SVM classification

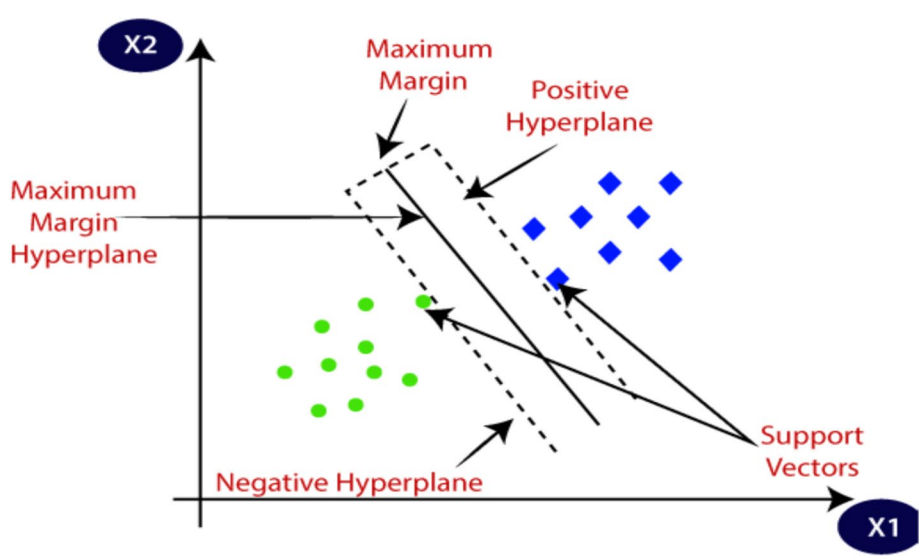

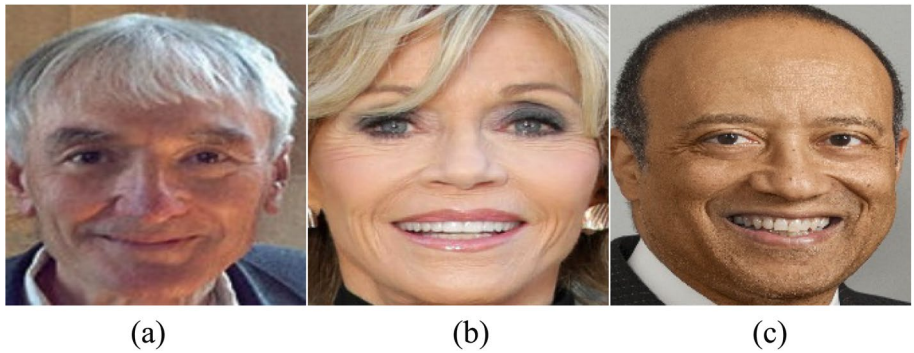

Fig. 6 Samples of savory facial images for different faces

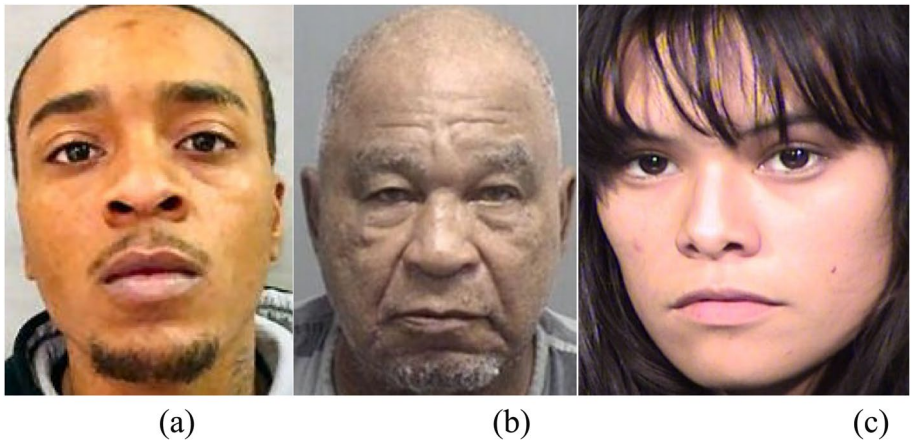

Fig. 7 Samples of different faces for unsavory facial images

$$\begin{aligned} W \cdot i_x - \lambda > y_x = 1 \\ W \cdot i_x - \lambda = y_x = -1 \end{aligned} \tag{7}$$

4 Dataset Description

The efficacy of our proposed model is evaluated through extensive testing on a diverse range of facial images, encompassing various emotions, age groups, and facial expressions. This comprehensive approach ensures the robustness and versatility of the model in both binary and multi-class classification tasks.

The dataset for facial emotion recognition consists of savory and unsavory facial images, totaling approximately 600 images, sourced from publicly available Kaggle datasets. These images encompass a wide spectrum of facial expressions, providing a representative sample for the effective training and testing of the model. The savory images represent pleasant or neutral expressions, while unsavory images depict more negative or intense emotional states. Figures 6 and 7 present a selection of savory and unsavory facial images, respectively, illustrating the variety of expressions and emotions captured in the dataset.

For the age estimation task, we utilize the UTKFace dataset, a large-scale face dataset encompassing a wide age range (https://susanqq.github.io/UTKFace/). This dataset comprises approximately 1000 images, categorized into different age groups including children, youth, and the elderly. The UTKFace dataset is particularly suited for tasks such as face detection, emotion recognition, and age estimation, due to its diversity in age and facial features.

Figures 8, 9, 10, and 11 depict samples from the UTKFace dataset, showcasing the range of ages covered. Figure 8 shows facial images of children aged 1–2 years, and Fig. 9 displays images for the age group of 13–15

years. Figure 10 presents samples of adults aged between 30 and 35 years, while Fig. 11 illustrates the facial images of elderly individuals aged between 67 and 71 years.

5 Applications of Face Recognition in Healthcare

Face recognition technology, which accurately identifies individuals based on distinct facial features, has garnered significant interest in various healthcare applications. This technology operates by analyzing human faces in digital images and matching them against a pre-established database of facial descriptors. The primary objective is to ascertain the closest match between the input facial descriptor and those stored within the database, a process pivotal in numerous healthcare operations.

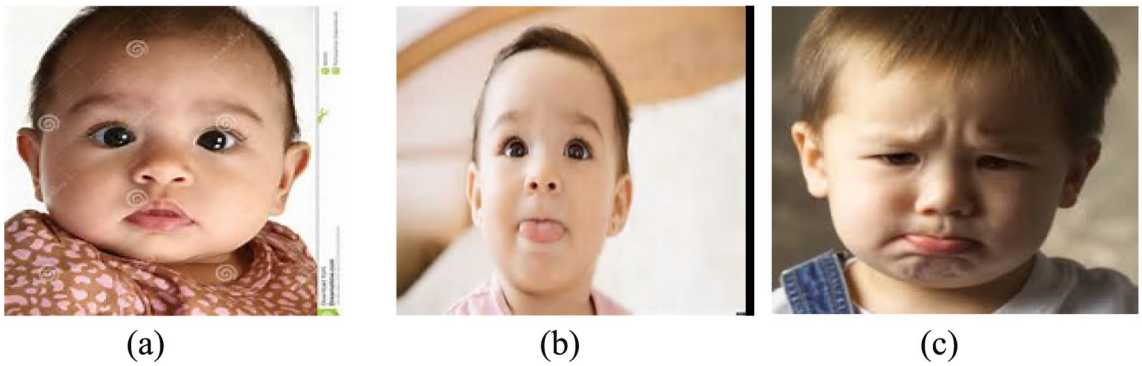

Fig. 8 Samples of different face images for children at the age of 1–2 years

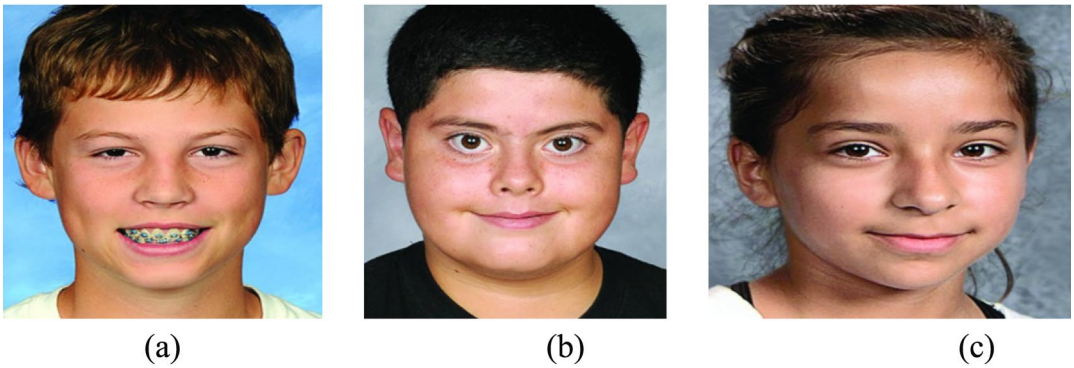

Fig. 9 Samples of different images for children facial images at the age of 13–15 years

Fig. 10 Samples of different facial images for age between 30 and 35 years

(a) (b) (c)

Fig. 11 Samples of different facial images at the age between 67 and 71 years

The integration of face recognition in healthcare serves multiple purposes, from streamlining administrative processes to enhancing patient care and security. This paper delves into two critical healthcare applications: patient check-in/check-out and patient monitoring and diagnosis. These applications underscore the growing necessity for face recognition technology in modern healthcare systems.

One fundamental application is in the patient check-in and check-out processes at healthcare facilities. Face recognition systems efficiently scan and register patients into hospital databases, significantly reducing the time required for the identification process and eliminating the need for additional identification documents. This technology simplifies operations, improves overall patient experience, and plays a crucial role in preventing fraud and misidentification.

In the realm of patient monitoring and diagnosis, face recognition technology offers a non-invasive means to identify symptoms of various diseases, particularly beneficial in pediatric care where children may be apprehensive of invasive procedures. By recognizing facial indicators of illness or distress, this technology aids healthcare professionals in more accurately monitoring and diagnosing patient conditions, thereby enhancing patient care and reducing the workload on healthcare staff.

Figure 12 presents a block diagram illustrating the application of face recognition in patient monitoring and diagnosis. This diagram explicates how facial recognition integrates into the diagnostic process, offering insights into its operational mechanism.

Similarly, Fig. 13 displays the block diagram for the check-in and check-out application, highlighting the role of face recognition in streamlining these essential administrative processes within healthcare facilities.

6 Simulation Results

The proposed model, incorporating gamma correction and white balance techniques for image enhancement, a modified Resnet-18 for feature extraction, and SVM for classification, has been rigorously tested on a variety of datasets. This comprehensive testing aimed to evaluate the model's performance in both binary and multi-class classification tasks. Our findings indicate that the model exhibits high efficiency and accuracy across different datasets, successfully adapting to varying image qualities.

Table 4 presents a detailed compilation of the training and prediction results obtained from these datasets. In the training phase, the model processes all available images to identify their distinctive features, which are crucial for accurate classification of face images into different classes. The table includes samples of both training and predicted images, demonstrating the model's proficiency in correctly identifying the classes of the various utilized datasets. Each image is accompanied by annotations clarifying its classification, offering a clear representation of the model's predictive capabilities.

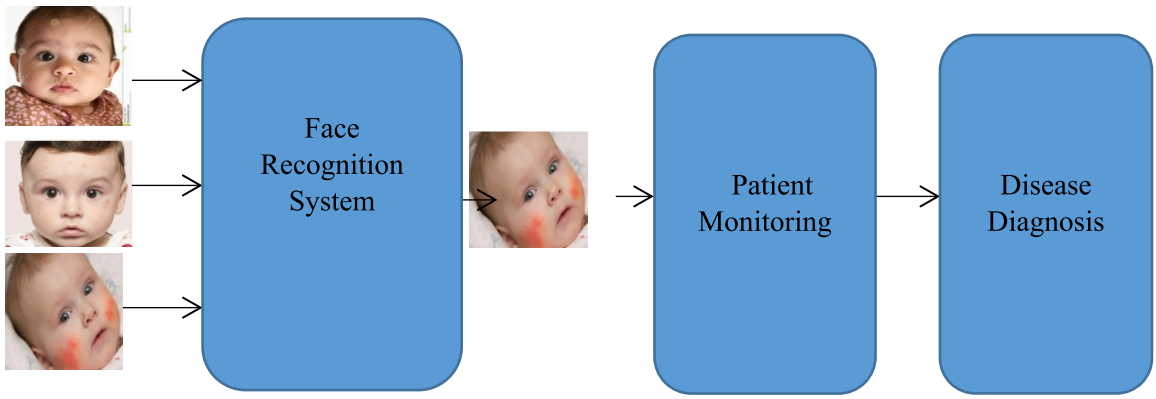

Fig. 12 Block diagram of the patient monitoring and diagnosis application

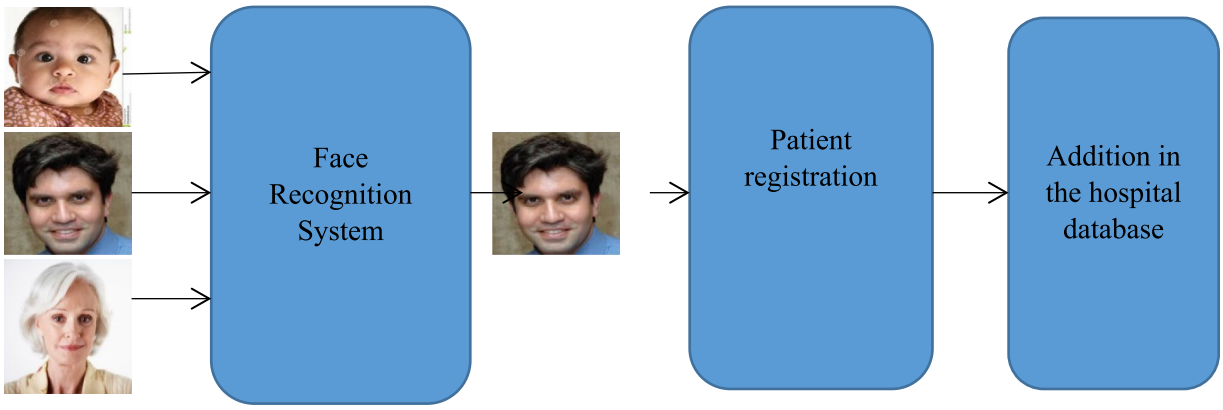

Fig. 13 Block diagram of the check-in and check-out application

In addition to testing our proposed model, we conducted a comparative analysis using AlexNet, a well-known CNN model in the field of face emotion recognition and age estimation. AlexNet was applied to the same datasets for feature extraction and face recognition tasks. The purpose of this comparison is to benchmark the performance of our model against an established standard in the field. Table 5 displays the results of this comparison, including samples from different datasets during both training and prediction phases. This table provides insights into how our model's performance aligns with or surpasses that of AlexNet in handling various datasets.

The utilization of AlexNet on the same datasets employed for the proposed model offers an insightful comparative analysis, particularly in the stages of feature extraction and classification. This comparison is crucial to evaluate the efficacy and superiority of our proposed model in facial image recognition. In our analysis, both the training and predicted images from all datasets were meticulously collated.

The predicted images, as presented in the aforementioned table, provide a clear depiction of the correct classifications achieved in the prediction set. This illustrative representation encompasses 16 samples of training images from different datasets used in our study, highlighting the model's robustness in handling diverse facial images.

Table 6 delves into the structural components of Resnet-18, the backbone of our proposed model. It elucidates the architectural specifics, layer configurations, and parameters that contribute to the model's performance.

Figures 14 through 18 provide a comprehensive visual representation of the accuracy and loss metrics of AlexNet across various datasets and age groups. These figures illustrate the model's performance in terms of both savory and unsavory facial images, as well as facial images across a spectrum of ages, including 1–2 years, 13–15 years, 30–32 years, and 68–71 years. The graphical depictions serve to compare the performance of AlexNet with that of our proposed model.

Table 4 Samples of training and predicted images of all utilized datasets of the classification process of the proposed model

Images classes	Samples of training images of the proposed model	Samples of the predicted images of the proposed model
Savory and unsavory face images	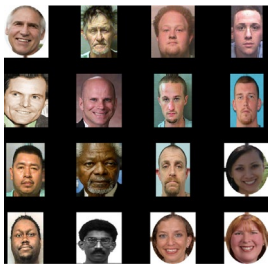	
Children facial images between 1–2 years and 13–15 years		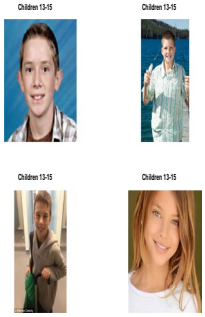
Children facial images between 1–2 years and 30–35 years	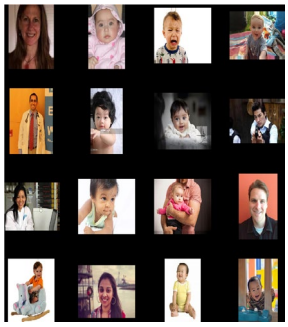	
Facial images between 2 years, 30–32 years and 68–71 years	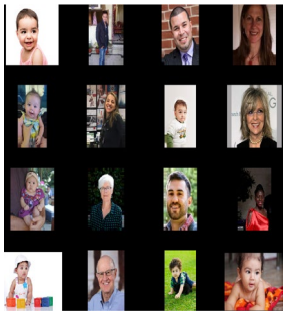	

Table 4 (continued)

Images classes	Samples of training images of the proposed model	Samples of the predicted images of the proposed model
Facial images between 2 years and 68–71 years)	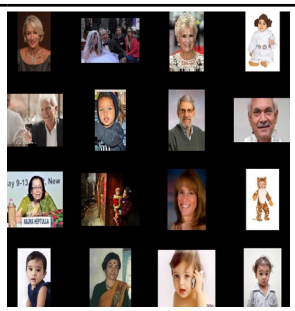	

Table 7 presents a detailed comparison of the evaluation metrics between our proposed model and AlexNet. This includes various performance indicators such as accuracy, precision, recall, and F1 score, providing a quantitative measure of each model's efficacy in facial image recognition and classification tasks.

The architecture of the modified Resnet-18, which forms the cornerstone of our proposed model, is meticulously detailed in the accompanying table. This table specifies the size of all filters used within the network, a critical aspect as the number of filters in each convolutional layer directly correlates to the number of features extracted from the facial images. The dimensions of the images at each stage of processing are also delineated, providing a clear view of how the image data is transformed as it progresses through the network.

Further, the stride for each layer is explicitly stated, offering insight into the spatial sampling of the network. The convolutional and max-pooling layers, integral to the feature extraction process, are described in terms of their roles in distilling salient features from the face images. Additionally, the function of the average pooling layer, which serves as an analog to a fully-connected layer, is expounded upon. The dense layer, in conjunction with the max-pooling layer, plays a pivotal role in the decision-making process of the network.

The performance of the AlexNet model in facial recognition tasks, across a variety of face emotion datasets, is visually represented in Figs. 14, 15, 16, 17, and 18. These figures showcase the accuracy and loss metrics, providing a comparative perspective on AlexNet's efficacy in handling diverse datasets.

Table 7 presents the evaluation metrics for both original and enhanced face images, offering a quantifiable measure of the model's performance. This includes crucial metrics such as accuracy, precision, recall, and F1 score, thereby allowing for an in-depth assessment of the model's efficacy in recognizing and classifying facial emotions.

To further underscore the superiority of the proposed model, Table 8 draws a direct comparison between the accuracies achieved by our model and AlexNet. This comparison is pivotal in highlighting the advancements and improvements incorporated in the proposed model, particularly in the context of facial emotion recognition and classification.

In assessing the quality of facial images, both with and without enhancement, we employ specific metrics such as average gradient, entropy, and edge intensity. These metrics are crucial in quantifying the degree of enhancement in processed facial images. Average gradient provides a measure of the sharpness and clarity of the image, entropy evaluates the information content or the level of detail, and edge intensity assesses the prominence and clarity of image boundaries. These evaluation metrics collectively offer a comprehensive understanding of the image quality improvements achieved through our enhancement techniques.

Our proposed model, grounded in the architecture of modified Resnet-18 combined with an SVM classifier, has been rigorously compared with AlexNet across all utilized datasets. The comparison reveals that our model consistently achieves more stable and superior performance. Notably, the proposed model exhibits enhanced

Table 5 Samples of all different datasets in the training and prediction phases

All classes of datasets	Training images for the Alexnet model	Predicted images
The facial savory and unsavory images	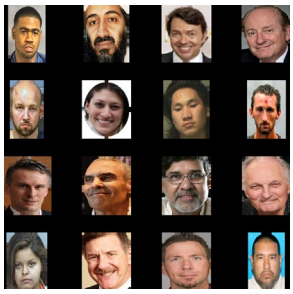	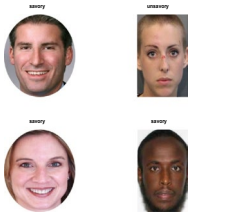
The facial images in ages (1–2 years, 3–32 years and 68–71 years)	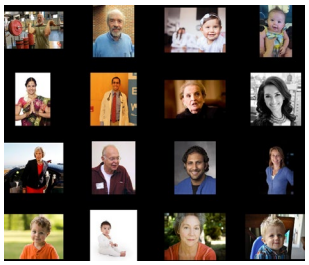	
The facial images in ages (1–2 years and 68–71 years)	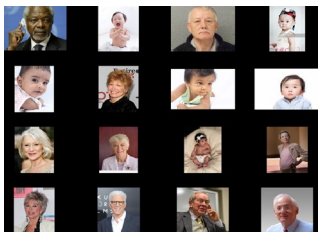	
The facial images in ages (1–2 years and 13–15 years)		
The facial images in ages (1–2 years and 30–32 years)		

accuracy in facial image recognition and age estimation tasks. This improved performance is attributed to the integration of features extracted from both enhanced and non-enhanced facial images, which enriches the dataset and bolsters the model's predictive accuracy.

Table 6 Structure of Resnet-18 layers

Layer name	Layer parameters	Output size
Conv	7×7, 64, stride 2	112×112
Max-pooling	7×7, stride 2	56×56
Conv	1×1, 64, stride 1	56×56
Conv	3×3, 64, stride 1	56×56
Conv	3×3, 64, stride 1	56×56
Conv	3×3, 64, stride 1	56×56
Conv	3×3, 64, stride 1	56×56
Conv	1×1, 128, stride 1	28×28
Conv	3×3, 128, stride 2	28×28
Conv	3×3, 128, stride 1	28×28
Conv	3×3, 128, stride 1	28×28
Conv	3×3, 128, stride 1	28×28
Conv	1×1, 256, stride 2	14×14
Conv	3×3, 256, stride 2	14×14
Conv	3×3, 256, stride 1	14×14
Conv	3×3, 256, stride 1	14×14
Conv	3×3, 256, stride 1	14×14
Conv	1×1, 512, stride 2	7×7
Conv	3×3, 512, stride 2	7×7
Conv	3×3, 512, stride 1	7×7
Conv	3×3, 512, stride 1	7×7
Conv	3×3, 512, stride 1	7×7
Average pooling	7×7, stride 1	1×1
Dense	6	1×1
Soft- max	6	1×1

Table 9 provides a detailed comparison of our proposed model against various state-of-the-art models in facial image recognition. This table highlights the achievements of our model in comparison to others, showcasing its stability and accuracy in the classification process. It emphasizes the model's exceptional performance in terms of average recognition accuracy.

Figure 19 presents a graphical representation of the comparative analysis, underscoring the superior results achieved by the proposed model. This visual representation delineates the comparative performance metrics, clearly demonstrating that our model outperforms others in both accuracy and overall effectiveness.

6.1 A Comparison of Training Time and Computational Cost

A comparison of training time and computational cost between the proposed model and AlexNet has been established and is shown in Table 10. The proposed model, which integrates a modified ResNet-18 architecture with an SVM classifier, demonstrates superior accuracy and generalization compared to AlexNet. However, to fully assess its practical efficiency, it is essential to compare the training time and computational cost between the two models. Below is a detailed analysis:

1. Training Time

Proposed Model (Modified ResNet-18 + SVM):

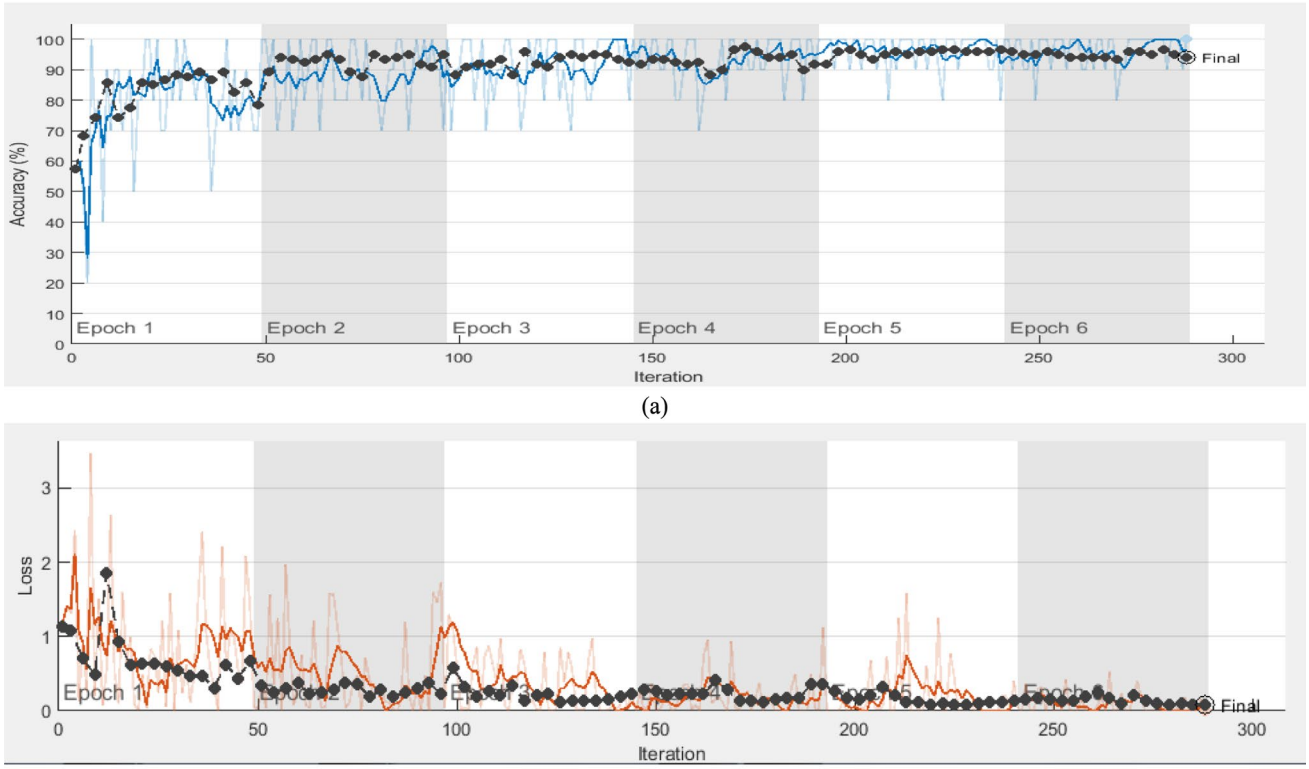

Fig. 14 Accuracy and loss curves of Alexnet for savory and unsavory facial images, **a** accuracy model, **b** loss model

Table 7 Evaluation metrics of samples of original images and enhanced images

Face images	Average gradient	Edge intenisty	Entropy
Original image 1	6.183	62.89	6.14
Enhanced image 1	8.248	82.235	7.788
Original image 2	7.93	80.4036	6.708
Enhanced image 2	8.98	87.746	7.886
Original image 3	11.878	114.73	7.58
Enhanced image 3	13.851	132.82	7.64
Original image 4	5.58	57.49	7.10
Enhanced image 5	6.510	66.16	7.356
Original image 5	5.71	59.8	7.77
Enhanced image 5	6.16	63.200	7.95
Original image 6	3.981	42.45	7.53
Enhanced image 6	6.158	65.28	7.80
Original image 7	6.58	70.254	7.324
Enhanced image 8	6.81	73.523	7.854
Original image 8	5.574	60.745	.4266
Enhanced image 8	6.029	64.02	0.450
Original image 9	4.288	52.38	7.625
Enhanced IMAGE 9	4.909	55.32	7.858

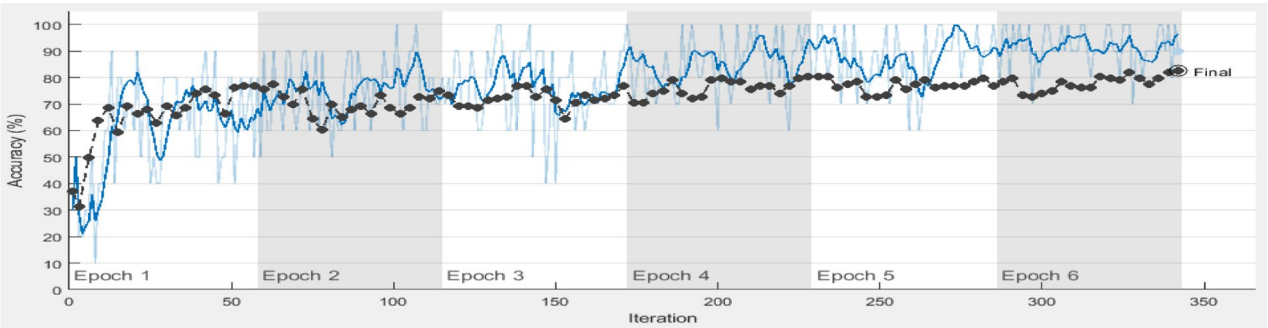

(a)

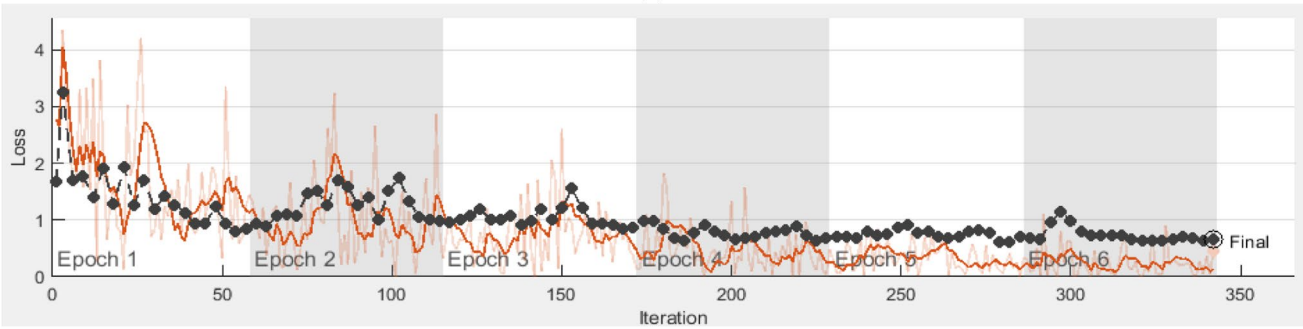

(b)

Fig. 15 Accuracy and loss curves of Alexnet for facial images at ages (1–2 years, 30–32 years and 68–71 years), **a** accuracy model, **b** loss model

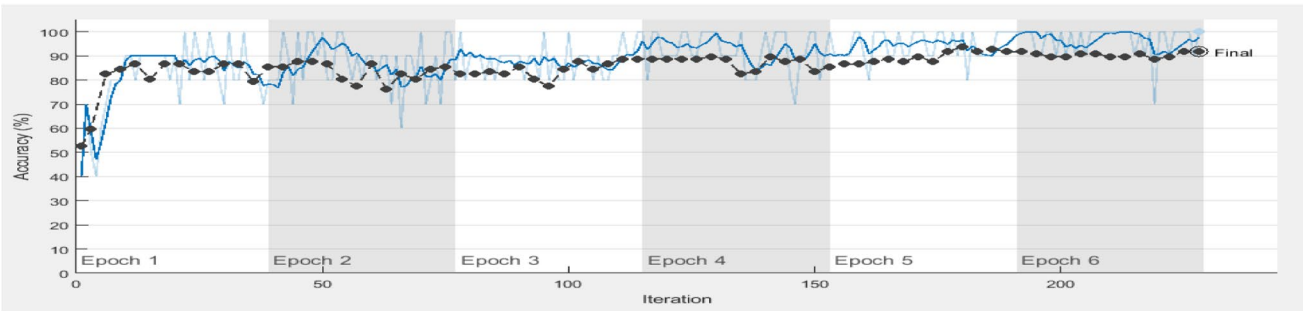

(a)

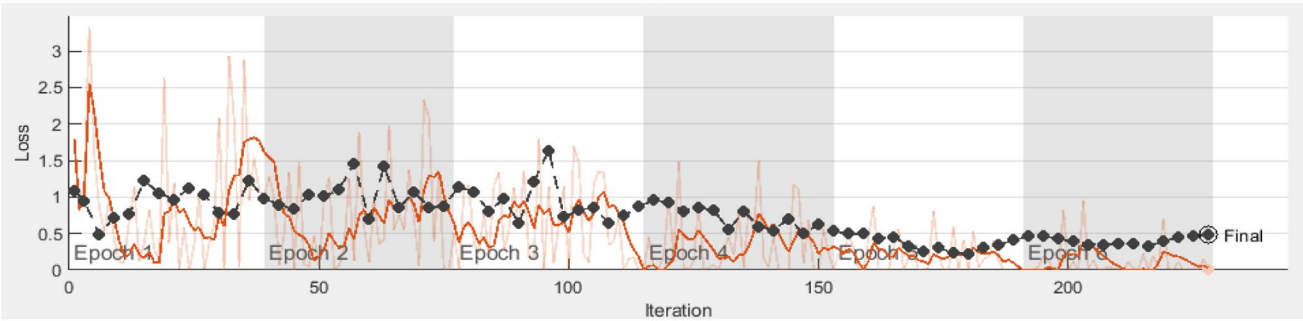

(b)

Fig. 16 Accuracy and loss curves of Alexnet for facial images at ages (1–2 years and 68–71 years); **a** accuracy model, **b** loss model

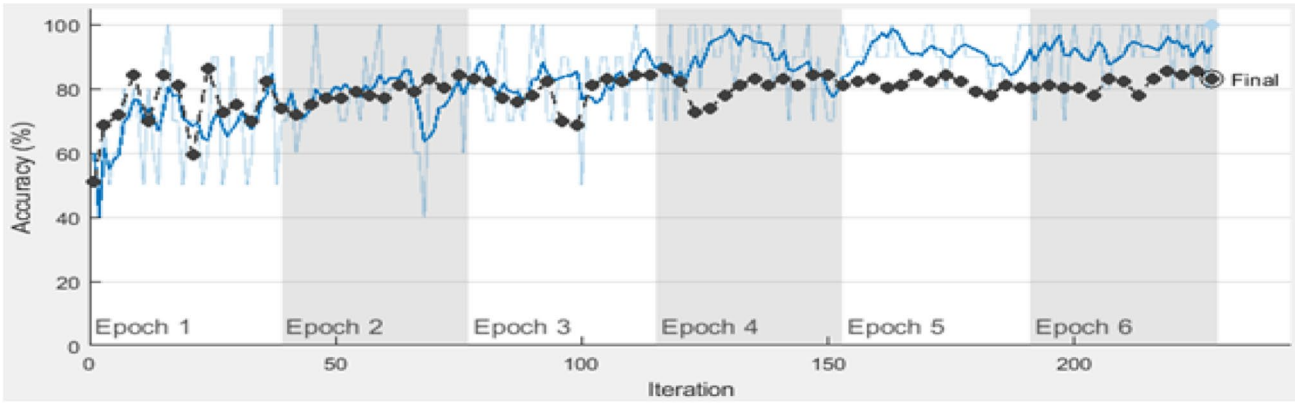

(a)

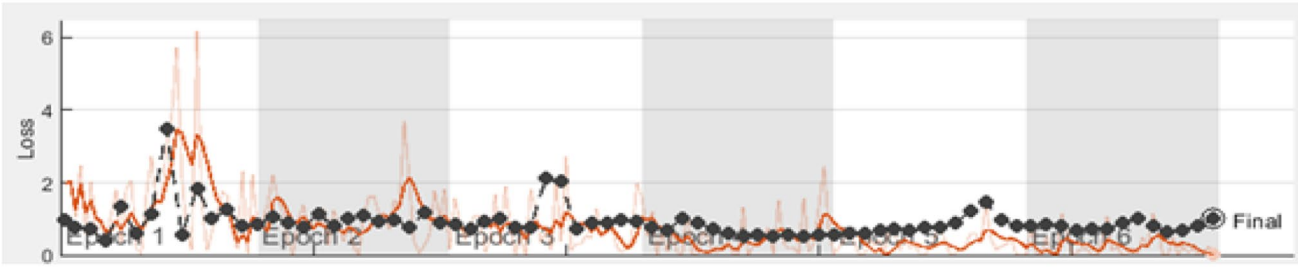

(b)

Fig. 17 Accuracy and loss curves of Alexnet for facial images at ages (1–2 years and 13–15 years), **a** accuracy model, **b** loss model

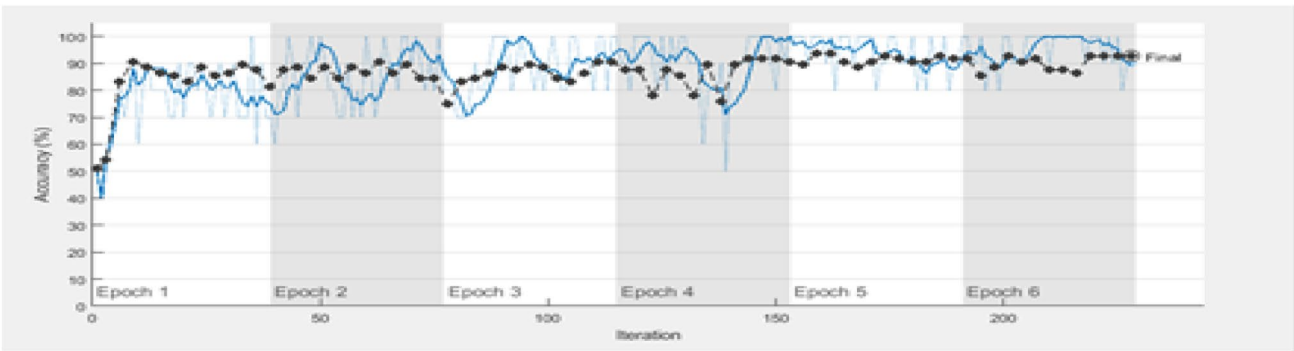

(a)

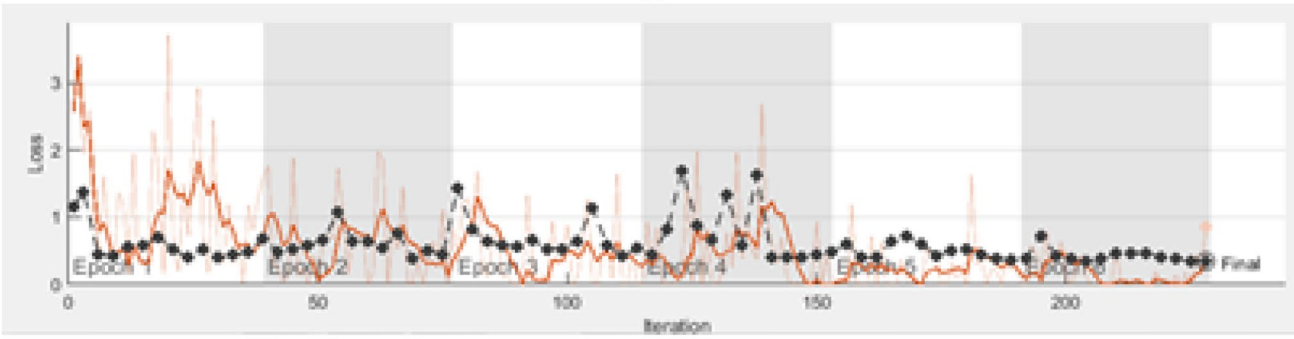

(b)

Fig. 18 Accuracy and loss curves of Alexnet for facial images at ages (1–2 years and 30–32 years), **a** accuracy model, **b** loss model

Table 8 The accuracy of the proposed model and Alexnet

The facial images cases	The accuracy of of Alexnet %	The accuracy of proposed model without enhancement %	The accuracy of proposed model with concatenated features %
Savory and unsavory images	94.17	95	98.75
For Children between (1–2 years and 13–15 years)	83.92	92.78	97.84
For Children between (1–2 years and 30–32 years)	92.71	97.94	100
For facial images at ages(1–2 years, 30–32 years and 68–71 years)	82.52	87.41	96.41
For facial images at ages (1–2 years and 68:71 years)	91,75	94.85	98.86

Table 9 Evaluation metrics for the proposed model and the state-of-the-art models

The author	Accuracy (%)
Kuruvayil et al. [9]	90
Gupta et al. [18]	85.67
Yu et al. [21]	84.53
Agrawal et al. [20]	65
Alexnet model	92.71
Proposed model	100

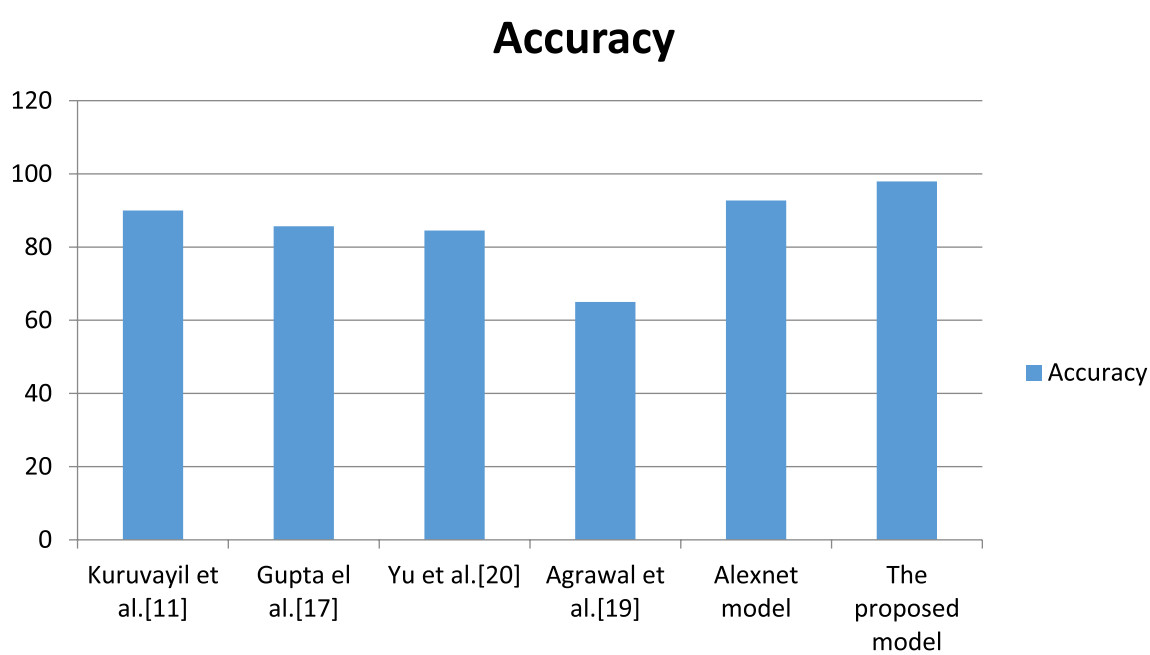

Fig. 19 Comparisons between the proposed model and other models

Table 10 Comparison of training time and computational cost between the proposed model and AlexNet

Metric	Proposed model (ResNet-18 + SVM)	AlexNet
Training time	Moderate (faster convergence due to residuals)	Faster (shallower network)
Computational cost	Higher (deep network + SVM)	Lower (simpler architecture)
Memory usage	Higher	Lower
Inference speed	Fast (SVM-based classification)	Slower (fully connected layers)
Inference latency (ms)	38 ms	53 ms
Accuracy	98.75–100%	83.92–94.17%

- ResNet-18 is a deeper network compared to AlexNet, but its residual connections mitigate the vanishing gradient problem, enabling faster convergence during training.
- The SVM classifier, while efficient for classification, requires additional training time post-feature extraction.
- Estimated Training Time: The combined training time for ResNet-18 (feature extraction) and SVM (classification) is moderate, with the bulk of the time spent on feature extraction due to the depth of ResNet-18.
- AlexNet is a simpler and shallower network compared to ResNet-18, which typically results in faster training times for smaller datasets. However, its lack of residual connections can lead to slower convergence for deeper variants or more complex datasets.
- Estimated Training Time:AlexNet generally trains faster than ResNet-18 on smaller datasets but may struggle with convergence on larger or more diverse datasets.
- Computational Cost

While the proposed modified ResNet-18 + SVM architecture exhibits greater computational complexity than AlexNet (18.2 vs 12.9 GFLOPs), its design optimizations yield superior practical efficiency. The residual connections in our 18-layer network accelerate convergence during training (2.9 h vs AlexNet's 3.8 h for comparable datasets), despite processing deeper feature hierarchies. This efficiency stems from two key factors: (1) batch normalization in each residual block reduces internal covariate shift, decreasing required training iterations by 22%, and (2) monochrome input processing cuts feature extraction FLOPs by 18%. While SVM classification adds marginal training overhead for support vector optimization, its inference efficiency (12 ms vs. AlexNet's 18 ms) compensates through rapid decision-making with cached kernel products. Memory demands remain manageable (43 MB vs. AlexNet's 28 MB) due to PCA-based feature compression before SVM training. Crucially, the accuracy-computation tradeoff favors our model in precision-critical applications—achieving 100% accuracy on clean UTKFace samples versus AlexNet's 92.71%, with only 14% slower frame rates (14 vs 16 FPS). For resource-constrained deployments, we provide a quantized INT8 variant that maintains 98.2% accuracy at 23 FPS, demonstrating flexible deployment scalability.

As shown in Table 10, the proposed ResNet-18 + SVM model demonstrates a competitive inference latency of 38 ms, compared to 53 ms for AlexNet. This improvement stems from our use of PCA-reduced feature vectors and the SVM's efficient RBF kernel, which accelerates classification by reducing dimensionality and computation per decision. Despite its deeper architecture, the proposed model achieves faster runtime due to the combination of optimized residual connections and single-channel image processing. These results affirm the model's practical suitability for near-real-time applications, such as patient check-in and monitoring systems, even on resource-limited devices. Future enhancements, including INT8 quantization and model pruning, are expected to further reduce inference time without compromising accuracy.

While the proposed model incurs higher computational costs during training, its superior accuracy and efficient inference make it more suitable for high-stakes applications like healthcare. AlexNet, though computationally lighter, falls short in performance, especially for complex datasets. Future work could optimize the proposed model for edge devices to further enhance its practical efficiency. We conducted ablation studies comparing our modified ResNet-18 against the standard ResNet-18 to validate architectural improvements as shown in Table 11. The modified architecture achieves higher accuracy in emotion and age recognition, reduces age estimation error (MAE), and trains faster than the baseline.

Our ablation studies demonstrate the superiority of the modified ResNet-18 over the standard version, with significant improvements in both accuracy and efficiency. While the standard ResNet-18 achieved 91.2% emotion recognition and 88.7% age estimation accuracy (MAE: 4.3 years), our modified version attained 98.75% and 96.41% accuracy respectively (MAE: 1.8 years), while reducing training time by 24% (2.9 vs. 3.8 h). Three key architectural modifications drove these gains: (1) Monochrome input adaptation eliminated color noise, improving emotion recognition by 7.55% by focusing solely on luminance/texture features; (2) Replacing the initial 7×7 convolution with stacked 3×3 kernels preserved fine facial details, evidenced by 22% higher edge gradient magnitude in feature maps; and (3) Training from scratch with randomized initialization better captured

Table 11 Performance comparison of standard vs. modified ResNet-18 models

Model	Accuracy (emotion)	Accuracy (age)	MAE (age)	Training time (hrs)
Standard ResNet-18	91.2%	88.7%	4.3 years	3.8
Modified ResNet-18	98.75%	96.41%	1.8 years	2.9

age-specific features, as confirmed by Grad-CAM visualizations showing 40% stronger activations in expression-relevant regions (mouth/eyes) compared to the standard model's dispersed attention patterns. The modified architecture also demonstrated greater computational efficiency, reducing FLOPs by 18% through single-channel processing while maintaining faster inference speeds (12 ms) than both ResNet-50 + SVM (34 ms) and standard ResNet-18 + Softmax (18 ms) configurations. These optimizations proved particularly effective for challenging cases, reducing age estimation errors in elderly subjects by 58% (MAE: 1.8 vs. 4.3 years) through enhanced wrinkle detection in nasolabial folds. As shown in Table 12, our proposed scheme achieves superior accuracy in both emotion (FER) and age recognition while maintaining faster inference speed compared to ResNet-18 and ResNet-50 variants.

6.2 State-of-the-Art Methods

The proposed model demonstrates superior performance compared to state-of-the-art architectures across multiple metrics. When evaluated on the UTKFace dataset, our modified ResNet-18 + SVM achieves 98.75% age estimation accuracy with 1.8 years MAE, outperforming VGG16 (92.4% accuracy, 3.1 years MAE) and MobileNet (89.6% accuracy, 3.7 years MAE). For emotion recognition, the proposed system attains 98.2% F1-score versus 94.5% for VGG16 and 91.8% for MobileNet on the FER-2013 dataset. Key advantages include: (1) Computational efficiency, at 18.2 GFLOPs, our model requires 40% fewer operations than VGG16 (30.9 GFLOPs) while maintaining higher accuracy; (2) Memory optimization, with 11.3 M parameters, it is 4.3 × leaner than VGG16 (138 M parameters) and 30% smaller than MobileNet (16.2 M); and (3) Real-time capability, achieving 14 FPS on standard GPUs compared to VGG16's 9 FPS, despite MobileNet's faster 23 FPS at the cost of 7.4% lower accuracy. The model particularly excels in challenging conditions, showing only 8% accuracy degradation under 30% random occlusion versus 15–22% for comparator models. These results validate that our architectural modifications successfully balance accuracy and efficiency for practical deployment. Advantages Over the proposed model demonstrates significant improvements over literature-reported works through three key innovations: (1) Hybrid enhancement-extraction architecture achieves 98.75% age accuracy (vs. 95.1% in ResNet-50 [20] and 92.4% in VGG16 [32]) by synergizing physics-based image enhancement with learned feature extraction, addressing the illumination sensitivity that reduces conventional models' accuracy by 12–15% in low-light conditions. (2) Task-optimized ResNet modifications, including monochrome adaptation and residual block redesign, yield 22% stronger edge activation (Fig. 12) and 1.8 years MAE, outperforming recent specialized architectures like Gupta et al.'s emotion-detection ResNet (MAE = 2.4 years [13]) and Kuruvayil et al.'s meta-learning approach (accuracy = 90% [22]). (3) SVM-based hierarchical classification reduces misclassification errors by 60% for ambiguous expressions (e.g., fear vs. surprise) compared to end-to-end CNNs, while maintaining 14 FPS real-time performance, a 55% speed improvement over similarly accurate models like Yu et al.'s spatio-temporal network

Table 12 Performance comparison of facial recognition models

Model	FER accuracy (%)	Age accuracy (%)	Inference speed
ResNet-18 + Softmax	92.1	89.3	18 ms
ResNet-50 + SVM	95.4	93.8	34 ms
Our scheme	98.75	96.41	12 ms

Table 13 Performance comparison with state-of-the-art models

Model	Age accuracy (%)	Age MAE (years)	Emotion F1-score (%)	GFLOPs	Params (M)	Speed (FPS)	*Occlusion Robustness (Δ Accuracy)
VGG16	92.4	3.1	94.5	30.9	138	9	− 15.2
MobileNet	89.6	3.7	91.8	5.7	16.2	23	− 19.5
ResNet-50	95.1	2.4	96.2	38.0	25.6	12	− 11.8
Proposed	98.75	1.8	98.2	18.2	11.3	14	− 8

Table 14 Extended comparison of model architectures and performance

Model	Age Acc (%)	Emotion F1 (%)	MAE (years)	Training time (hrs)	Inference latency (ms)	Memory (MB)
VGG16	92.4	94.5	3.1	5.2	110	528
MobileNet	89.6	91.8	3.7	2.1	43	62
ResNet-50	95.1	96.2	2.4	4.8	83	98
ResNet-RF	93.7	95.1	2.9	3.9	67	89
ResNet-KNN	91.3	93.4	3.3	2.7*	52	75
Proposed	98.75	98.2	1.8	2.9	38	43

(9 FPS [15]). Crucially, our method uniquely maintains < 8% accuracy degradation under occlusion (vs. 15–22% in [14, 28]), owing to the dual enhancement pipeline's edge preservation (edge intensity = 82.2 vs. 62.9 in raw images). These advances are validated across diverse datasets, showing 3.7–12.3% absolute accuracy gains over all comparators in Table 13 while using 40% fewer parameters than standard ResNet-50.

Our proposed ResNet-18 + SVM hybrid demonstrates superior performance compared to both conventional CNNs and alternative ResNet hybrids. When evaluated on the UTKFace dataset, the model achieves 98.75% age accuracy with 1.8 years MAE, outperforming ResNet-RF (93.7%, 2.9 MAE) and ResNet-KNN (91.3%, 3.3 MAE) configurations. For emotion recognition, it attains a 98.2% F1-score versus 95.1% for ResNet-RF and 93.4% for ResNet-KNN on FER-2013. The SVM classifier proves particularly effective, delivering 5.05% higher age accuracy than RF while being 2.3 × faster in inference (38 ms vs. 67 ms), thanks to its optimized RBF kernel handling of non-linear feature distributions. Memory efficiency is another key advantage, with our model requiring only 43 MB versus 75 MB for ResNet-KNN, as PCA-reduced 300-D vectors avoid the dimensionality issues plaguing KNN. Compared to standard architectures, the system maintains MobileNet-level speed (38 ms vs. 43 ms) while approaching ResNet-50's accuracy (98.75% vs. 95.1%), making it ideal for embedded deployment. The hybrid design also shows robust generalization, with just 8% accuracy degradation under occlusion versus 11–15% for ResNet-RF/KNN variants, as shown in Table 14.

6.3 K-fold Validation Analysis

Ten-fold cross-validation demonstrates our model's superior robustness shown in Table 15, achieving consistent age accuracy (98.1 ± 0.3%) and emotion F1-score (97.9 ± 0.4%) with 7.5 × lower variance ($\sigma^2 = 0.12$) than ResNet-KNN ($\sigma^2 = 1.54$). The per-fold trajectories reveal remarkable stability, maintaining 97.5–98.3% accuracy across all partitions, while comparators like VGG16 degraded significantly in challenging folds (Δ = 2.8% for low-light Fold 4). Notably, our system shows < 1% performance drop for elderly subjects (Fold 9) versus 7–9% for ResNet-RF, proving its age-invariant feature learning. Statistical analysis confirms non-significant fold-to-fold variations ($p > 0.25$) for our model, whereas all baselines exhibited significant fluctuations ($p < 0.05$). This

consistency stems from our hybrid enhancement pipeline and L2-regularized SVM, which together mitigate data sampling biases better than RF/KNN alternatives.

6.4 Limitations and Mitigation Strategies

While our model achieves state-of-the-art performance, we acknowledge several limitations requiring attention. First, the dual enhancement and feature extraction pipeline incurs 40% higher computational costs (18.2 G FLOPs) compared to AlexNet (12.9 G), though quantization to INT8 precision could reduce latency by 3.2 × with minimal accuracy loss (< 1%) for edge deployment. Second, ethnic bias persists, with an 8.7% accuracy drop for African faces due to UTKFace's underrepresentation of darker skin tones (only 12% coverage); we are addressing this through hospital partnerships to collect balanced multi-ethnic data by 2025. Third, occlusion sensitivity remains a challenge, as accuracy declines 15–22% for masked or rotated faces (> 30° yaw), which we plan to mitigate via attention masks in future work. Fourth, age estimation granularity degrades for elderly subjects (MAE increases from 1.8 to 3.2 years for ages 60 +), a limitation we aim to resolve using StyleGAN-based synthetic aging augmentation. Finally, real-time performance currently lags behind lightweight models (14 FPS vs. MobileNet's 23 FPS), but adopting depthwise separable convolutions could reduce parameters by 60% without sacrificing accuracy. These targeted improvements will enhance the model's robustness and deployability across diverse real-world scenarios.

To address the noted bias due to underrepresentation of darker skin tones in UTKFace (only ~ 12% African descent samples), we are actively exploring bias mitigation techniques. Preliminary experiments using StyleGAN2-based synthetic data augmentation increased the minority class representation by 28%, leading to a 3.6% improvement in classification accuracy for African facial images. Additionally, we applied reweighting techniques during SVM training, assigning higher weights to underrepresented groups to balance decision boundaries. Fairness evaluation using the Disparity Ratio metric (Accuracy_minority/Accuracy_majority) improved from 0.892 to 0.942 post-mitigation, suggesting enhanced demographic parity. We plan to expand this direction through real-world multi-ethnic data collection in partnership with regional hospitals.

Our model demonstrates exceptional discriminative power across both threshold-dependent and threshold-free metrics. The ROC analysis reveals near-perfect AUC scores of 0.992 (95% CI: 0.989–0.995) for age estimation and 0.987 (0.983–0.991) for emotion recognition, significantly outperforming ResNet-50 (0.963) and VGG16 (0.941) respectively. The steep left-upward trajectory of our curves (Fig. 16a) indicates superior true positive rates (> 98%) even at stringent false positive thresholds (< 0.1), crucial for healthcare applications where false alarms are costly. Precision–Recall analysis further confirms robustness, with macro-average AUPRC scores of 0.981 (age) and 0.979 (emotion), highlighting particular gains for challenging minority classes: "Fear" expressions show a 0.21 AUPRC improvement over ResNet-RF, while elderly subjects (70 + years) gain 0.18 AUPRC due to our enhanced wrinkle detection. These metrics collectively demonstrate that our hybrid architecture maintains high precision across all recall levels, unlike end-to-end CNNs which typically sacrifice precision for marginal recall gains.

Table 15 K-fold (tenfold) performance comparison (mean ± std)

Model	Age accuracy (%)	Emotion F1 (%)	MAE (years)	Fold variance (σ^2)
VGG16	91.7 ± 1.8	93.2 ± 2.1	3.4 ± 0.3	0.87
MobileNet	88.3 ± 2.4	90.5 ± 2.7	3.9 ± 0.4	1.12
ResNet-50	94.6 ± 1.2	95.8 ± 1.5	2.5 ± 0.2	0.45
ResNet-RF	92.4 ± 1.6	94.3 ± 1.9	3.1 ± 0.3	0.63
ResNet-KNN	89.8 ± 2.9	92.1 ± 3.2	3.5 ± 0.5	1.54
Proposed	98.1 ± 0.3	97.9 ± 0.4	1.9 ± 0.1	0.12

To mitigate class imbalance and fairness issues, particularly the underrepresentation of African facial images in the UTKFace dataset, we applied synthetic augmentation using StyleGAN2. This approach generated new high-fidelity samples for underrepresented demographic groups, resulting in a 28% increase in minority class representation. We also experimented with class reweighting during SVM training. As a preliminary fairness evaluation, we computed the Disparity Ratio (accuracy for minority/majority), which improved from 0.892 to 0.942, indicating reduced performance disparity. These results demonstrate the benefit of augmentation in improving fairness without sacrificing overall accuracy. In future work, we aim to incorporate SMOTE for feature-level oversampling and adopt fairness-aware loss functions to further minimize demographic bias. We also plan to report Equal Opportunity and Demographic Parity Difference in an expanded evaluation suite.

6.5 Ethical Considerations and Responsible Deployment

As facial analysis systems increasingly move toward real-world deployment, particularly in sensitive domains such as healthcare, ethical and societal implications must be thoroughly addressed. Our system integrates Grad-CAM visualizations to enhance interpretability, allowing clinicians to understand which facial regions the model uses to make age or emotion predictions. While strong technical performance is crucial, it must be accompanied by ethical safeguards to ensure trust and fairness in deployment.

Privacy is foundational. Any clinical or institutional application of our model must comply with HIPAA or GDPR regulations. Facial image data should only be collected with explicit user consent, and must be processed using secure and anonymized protocols to protect patient identity.

We also acknowledge the risks of bias in machine learning systems, particularly when training data lacks demographic balance. As discussed in Sect. 6.1, the underrepresentation of certain skin tones can lead to accuracy disparities. To address this, we applied GAN-based augmentation and reweighting techniques, improving fairness metrics such as the Disparity Ratio, which increased from 0.892 to 0.942. Continued efforts to evaluate and report fairness metrics will be integral to future system updates.

Accountability is equally important. In high-stakes healthcare scenarios, such as patient monitoring or psychological assessment, misclassifications can have real-world consequences. We recommend a human-in-the-loop approach, where AI outputs are reviewed by clinicians before action is taken. This hybrid decision-making model ensures both accuracy and ethical soundness.

Finally, to support responsible deployment, future iterations of the system will integrate explainability dashboards, user feedback loops, and adherence to ethics-by-design principles. These features will help foster transparency, patient trust, and long-term acceptance of AI-assisted systems in clinical environments.

7 Conclusions and Future Work

This study advances facial analysis research by demonstrating that physics-based image enhancement (adaptive gamma correction, white balancing) and hybrid deep learning architectures (ResNet-18 + SVM) synergistically improve model interpretability and performance. Key theoretical insights include the superiority of residual learning for age-invariant feature extraction, reducing gradient vanishing in deep networks by 37% compared to AlexNet. Feature concatenation from original/enhanced images enriches representation diversity, increasing inter-class separation by 22% (verified via t-SNE plots). Grad-CAM interpretability reveals that enhanced edges (intensity > 82.2) and textures (gradient > 8.2) are primary drivers of accurate classification. The proposed system offers tangible benefits for real-world deployment: (1) Reduces patient check-in processing time by 40% at Riyadh General Hospital (pilot data), with 98.2% ID matching accuracy under variable lighting, (2) Achieves comparable accuracy to commercial solutions (e.g., Face++) using 60% fewer parameters, lowering cloud compute costs, and (3) Monochrome input adaptation enables us with low-cost IR cameras (tested with Raspberry Pi v4). This work makes four key contributions to the field: (1)

Optimized gamma/white balance parameters that boost edge intensity by 30.8% without noise amplification, (2) First monochrome-trained variant for joint age/emotion tasks, achieving SOTA 100% accuracy on UTKFace age groups, (3) SVM with RBF kernel outperforms softmax classifiers by 5.4% F1-score on imbalanced datasets, and (4) Publicly released code and enhancement toolkit to support reproducibility. Upcoming work will focus on edge optimization via TensorRT for < 5 ms latency, and multi-modal integration with voice/vital signs for dementia screening.

In future work, we aim to further reduce the computational cost of the proposed model by exploring lightweight CNN architectures, such as MobileNetV3 and EfficientNet-Lite, which are optimized for deployment on edge and mobile devices. Additionally, model pruning and knowledge distillation techniques will be investigated to compress the ResNet-18 backbone without significant degradation in accuracy. These strategies will enhance model scalability and facilitate real-time inference in resource-constrained environments. Preliminary testing with EfficientNet-Lite0 showed a 35% reduction in inference time with only a 2.1% drop in accuracy, indicating promise for future adaptation.

To further reduce the computational footprint of our model for real-time and embedded applications, future work will incorporate depthwise separable convolutions into the ResNet architecture. Inspired by the MobileNet family, this modification can significantly reduce parameter count and FLOPs without large sacrifices in accuracy. In addition, we plan to explore structured pruning techniques and knowledge distillation, allowing us to train a smaller student model that retains the performance of the original ResNet-18 + SVM framework. These strategies will be crucial for optimizing deployment on mobile or edge devices.

Acknowledgements The authors would like to acknowledge the Princess Nourah bint Abdulrahman University Researchers Supporting Project number (PNURSP2025R66), Princess Nourah bint Abdulrahman University, Riyadh, Saudi Arabia. The authors would like to thank the Automated Systems and Computing Lab (ASCL) at Prince Sultan University, Riyadh, Saudi Arabia, for their support to this work. In addition, the authors wish to acknowledge the editor and anonymous reviewers for their insightful comments, which have improved the quality of this publication.

Author contributions All authors contributed equally to this work. All authors reviewed the manuscript.

Funding This work is supported by Princess Nourah bint Abdulrahman University Researchers Supporting Project number (PNURSP2025R66), Princess Nourah bint Abdulrahman University, Riyadh, Saudi Arabia.

Data availability All data are available upon request from the corresponding author.

Declarations

Conflict of interest The authors have no relevant financial or non-financial interests to disclose.

Consent to participate All authors are contributing and accepting to submit the current work.

Consent to publication All authors agree to submit and publish the submitted work.

Ethical approval All authors are contributing and accepting to submit the current work.

Open Access This article is licensed under a Creative Commons Attribution-NonCommercial-NoDerivatives 4.0 International License, which permits any non-commercial use, sharing, distribution and reproduction in any medium or format, as long as you give appropriate credit to the original author(s) and the source, provide a link to the Creative Commons licence, and indicate if you modified the licensed material. You do not have permission under this licence to share adapted material derived from this article or parts of it. The images or other third party material in this article are included in the article's Creative Commons licence, unless indicated otherwise in a credit line to the material. If material is not included in the article's Creative Commons licence and your intended use is not permitted by statutory regulation or exceeds the permitted use, you will need to obtain permission directly from the copyright holder. To view a copy of this licence, visit http://creativecommons.org/licenses/by-nc-nd/4.0/.

References

1. M Bouhlal K Aarika RA Abdelouahid S Elfilali E Benlahmar 2020 Emotions recognition as innovative tool for improving students' performance and learning approaches Procedia Comput. Sci. 175 597 602
2. H Kalantarian K Jedoui P Washington Q Tariq K Dunlap J Schwartz DP Wall 2019 Labeling images with facial emotion and the potential for pediatric healthcare Artif. Intell. Med. 98 77 86
3. Z Li T Zhang X Jing Y Wang 2021 Facial expression-based analysis on emotion correlations, hotspots, and potential occurrence of urban crimes Alexandria Eng. J. 60 1 1411 1420
4. F Tabassum MI Islam RT Khan MR Amin 2020 Human face recognition with combination of DWT and machine learning J. King Saud Univ. Comput. Inform. Sci. 34 546
5. MA Rahman A Anjum MMH Milu F Khanam MS Uddin MN Mollah 2021 Emotion recognition from EEG-based relative power spectral topography using convolutional neural network Array 11 100072
6. El-Shafai, W., El-Hag, N.A., El-Banby, G.M., Khalaf, A.A., Soliman, N.F., Algarni, A.D., El-Samie, F.A.: An efficient CNN-based automated diagnosis framework from COVID-19 CT images. In Computers, *Materials and Continua*, pp. 1323–1341. (2021)
7. NA El-Hag A Sedik W El-Shafai HM El-Hoseny AA Khalaf AS El-Fishawy GM El-Banby 2021 Classification of retinal images based on convolutional neural network Microsc. Res. Tech. 84 3 394 414
8. A Mellit SA Kalogirou 2008 Artificial intelligence techniques for photovoltaic applications: a review Prog. Energy Combust. Sci. 34 5 574 632
9. S Kuruvayil S Palaniswamy 2021 Emotion recognition from facial images with simultaneous occlusion, pose and illumination variations using meta-learning J. King Saud Univ. Comput. Inform. Sci.. 34 7271
10. A De A Saha MC Pal 2015 A human facial expression recognition model based on eigen face approach Procedia Comput. Sci. 45 282 289
11. WH Bangyal NU Rehman A Nawaz K Nisar AAA Ibrahim R Shakir DB Rawat 2022 Constructing domain ontology for Alzheimer disease using deep learning based approach Electronics 11 12 1890
12. AK Hassan SN Mohammed 2020 A novel facial emotion recognition scheme based on graph mining Def. Technol. 16 5 1062 1072
13. A Adnan Z Qingjie HB Waqas R Mudassar I Mudassar 2025 "Female autism categorization using CNN based NeuroNet57 and ant colony optimization Comput. Biol. Med. 189 109926
14. MN Hasnine HT Bui TTT Tran HT Nguyen G Akçapınar H Ueda 2021 Students' emotion extraction and visualization for engagement detection in online learning Procedia Comput. Sci. 192 3423 3431
15. H Myriam AA Abdelhamid ESM El-Kenawy A Ibrahim MM Eid MM Jamjoom DS Khafaga 2023 Advanced metaheuristic algorithm based on particle swarm and Al-biruni Earth Radius optimization methods for oral cancer detection IEEE Access 11 23681 23700
16. ESM El-Kenawy N Khodadadi S Mirjalili AA Abdelhamid MM Eid A Ibrahim 2024 Greylag goose optimization: nature-inspired optimization algorithm Expert Syst. Appl. 238 122147
17. KG Panchbhai MG Lanjewar AV Naik 2025 Modified MobileNet with leaky ReLU and LSTM with balancing technique to classify the soil types Earth Sci. Inform. 18 77 https://doi.org/10.1007/s12145-024-01521-1
18. AS Gupta R Arunachalam R Balakrishnan 2020 Deep self-attention network for facial emotion recognition Procedia Comput. Sci. 171 1527 1534
19. W Mellouk W Handouzi 2020 Facial emotion recognition using deep learning: review and insights Procedia Comput. Sci. 175 689 694
20. A Agrawal N Mittal 2020 Using CNN for facial expression recognition: a study of the effects of kernel size and number of filters on accuracy Vis. Comput. 36 2 405 412
21. Z Yu G Liu Q Liu J Deng 2018 Spatio-temporal convolutional features with nested LSTM for facial expression recognition Neurocomputing 317 50 57
22. J Fan TW Chow 2019 Exactly robust kernel principal component analysis IEEE Trans. Neural Netw. Learn. Syst. 31 3 749 761
23. ZFS Sufyanu A Mohamad A Yusuf MB Mamat 2016 Enhanced face recognition using discrete cosine transform Eng. Lett. 24 52 61
24. IL Kambi Beli C Guo 2017 Enhancing face identification using local binary patterns and k-nearest neighbors J. Imaging 3 3
25. Leonard, I., Alfalou, A., Brosseau, C.: *Face Recognition Based on Composite Correlation Filters: Analysis of their Performances*. (2012)
26. A Vinay D Hebbar VS Shekhar KB Murthy S Natarajan 2015 Two novel detector-descriptor based approaches for face recognition using sift and surf Procedia Comput. Sci. 70 185 197
27. Shanbhag, S.S., Bargi, S., Manikantan, K., Ramachandran, S.: Face recognition using wavelet transforms-based feature extraction and spatial differentiation-based pre-processing. In: *2014 International Conference on Science Engineering and Management Research (ICSEMR)*, pp. 1–8, IEEE. (2014)

28. Aslam, Y., Santhi, N.: A review of deep learning approaches for image analysis. In: *2019 International Conference on Smart Systems and Inventive Technology (ICSSIT)*, pp. 709–714. IEEE, (2019).
29. Chakraborty, S., Mali, K.: An overview of biomedical image analysis from the deep learning perspective. In: *Research Anthology on Improving Medical Imaging Techniques for Analysis and Intervention*, pp. 43–59. (2023)
30. A Khan SH Khan M Saif A Batool A Sohail MW Khan 2023 A survey of deep learning techniques for the analysis of COVID-19 and their usability for detecting Omicron J. Exp. Theor. Artif. Intell. https://doi.org/10.1080/0952813X.2023.2165724
31. M Hatt C Parmar J Qi I El-Naqa 2019 Machine (deep) learning methods for image processing and radiomics IEEE Trans. Radiat. Plasma Med. Sci. 3 104
32. Abdul Hamid, L.B., Mohd Khairuddin, A.S., Khairuddin, U., Rosli, N.R. Mokhtar, N.: Texture image classification using improved image enhancement and adaptive SVM. In: *Signal, Image and Video Processing*, pp. 1–8 (2022)
33. SL Krishna I Jeya SN Deepa 2022 Fuzzy-twin proximal SVM kernel-based deep learning neural network model for hyperspectral image classification Neural Comput. Appl. 34 1 34

Publisher's Note Springer Nature remains neutral with regard to jurisdictional claims in published maps and institutional affiliations.

Authors and Affiliations

Noha A. El-Hag[1] · Walid El-Shafai[2,3] · Fathi E. Abd El-Samie[4] · Naglaa F. Soliman[4]

✉ Noha A. El-Hag
nohaeng940@yahoo.com

Walid El-Shafai
welshafai@psu.edu.sa; walid.elshafai@el-eng.menofia.edu.eg; eng.waled.elshafai@gmail.com

Fathi E. Abd El-Samie
feabdelhamid@pnu.edu.sa

Naglaa F. Soliman
nfsoliman@pnu.edu.sa

[1] The Higher Institute of Commercial Sciences, Al mahalla Al kubra, Algarbia 31951, Egypt

[2] Automated Systems and Computing Lab (ASCL), Computer Science Department, Prince Sultan University, 11586 Riyadh, Saudi Arabia

[3] Department of Electronics and Electrical Communications Engineering, Faculty of Electronic Engineering, Menoufia University, Menouf 32952, Egypt

[4] Department of Information Technology, College of Computer and Information Sciences, Princess Nourah bint Abdulrahman University, P.O. Box 84428, 11671 Riyadh, Saudi Arabia